

Red-black tree

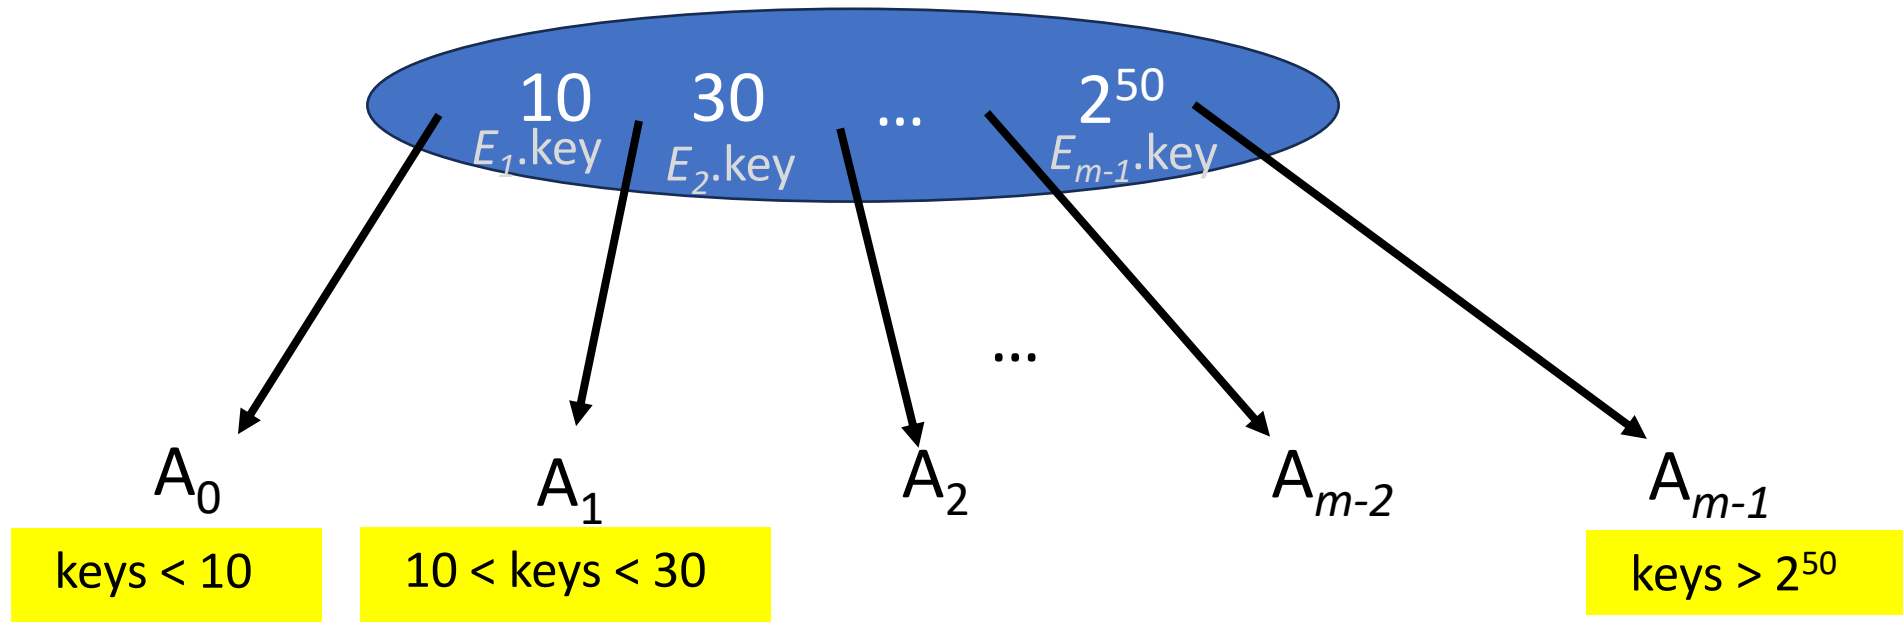
Ch. 10.3

Recall: m-way search tree

- Ch 10.2 AVL tree
- Ch 11.2 B-tree
 - 2-3 trees (B-tree of order 3)
 - 2-3-4 tree (B-tree of order 4)
- Ch 10.3 Red-black tree (An extension of 2-3-4 trees)
- Ch 11.3 B⁺-tree

We are here

Recall: m-way search tree




- If m is large, the node size becomes large and thus the cost of searching within a node increases.
- Can we just store one data pair in a node? \rightarrow Binary tree

Red-black tree is the binary tree simulating 4-way search tree (2-3-4 tree).

Recall: 2-3-4 tree

- The **nodes** store 1, 2 or 3 data pairs and have **2, 3 or 4 children**, respectively.
- All **external nodes** are at the **same level**.




Height = $O(\log n)$

Recall: 2-3-4 tree nodes


- 2-node

- Same as a binary node




- 3-node

- 2 data pairs, 3 children



- 4-node

- 3 data pairs, 4 children



Recall: Bottom-up insertion


In 2-3-4 tree, if a data pair is inserted into a 4-node:

- The node will be split and the mid data pair will be inserted into to the parent node.
- If the parent is also a 4-node, the splitting process will be repeated. (bottom-up splitting)

Red-Black tree

A binary search tree with the following properties:

- Edges are colored **red** or **black**.
- ***No two consecutive red edges*** on any root-external node path
- ***Same number of black edges*** on any root-external node path
(=*black height* of the tree)
- ***Edges connecting to external nodes are black.***



Relate 2-3-4 tree to red-black tree

- 2 node
 - One data pair


2-3-4



Red-Black




- 3-node
 - Two data pairs



Nodes connected by **red** edges represent the data pairs in a single node of 2-3-4 tree.

- 4-node
 - Three data pairs

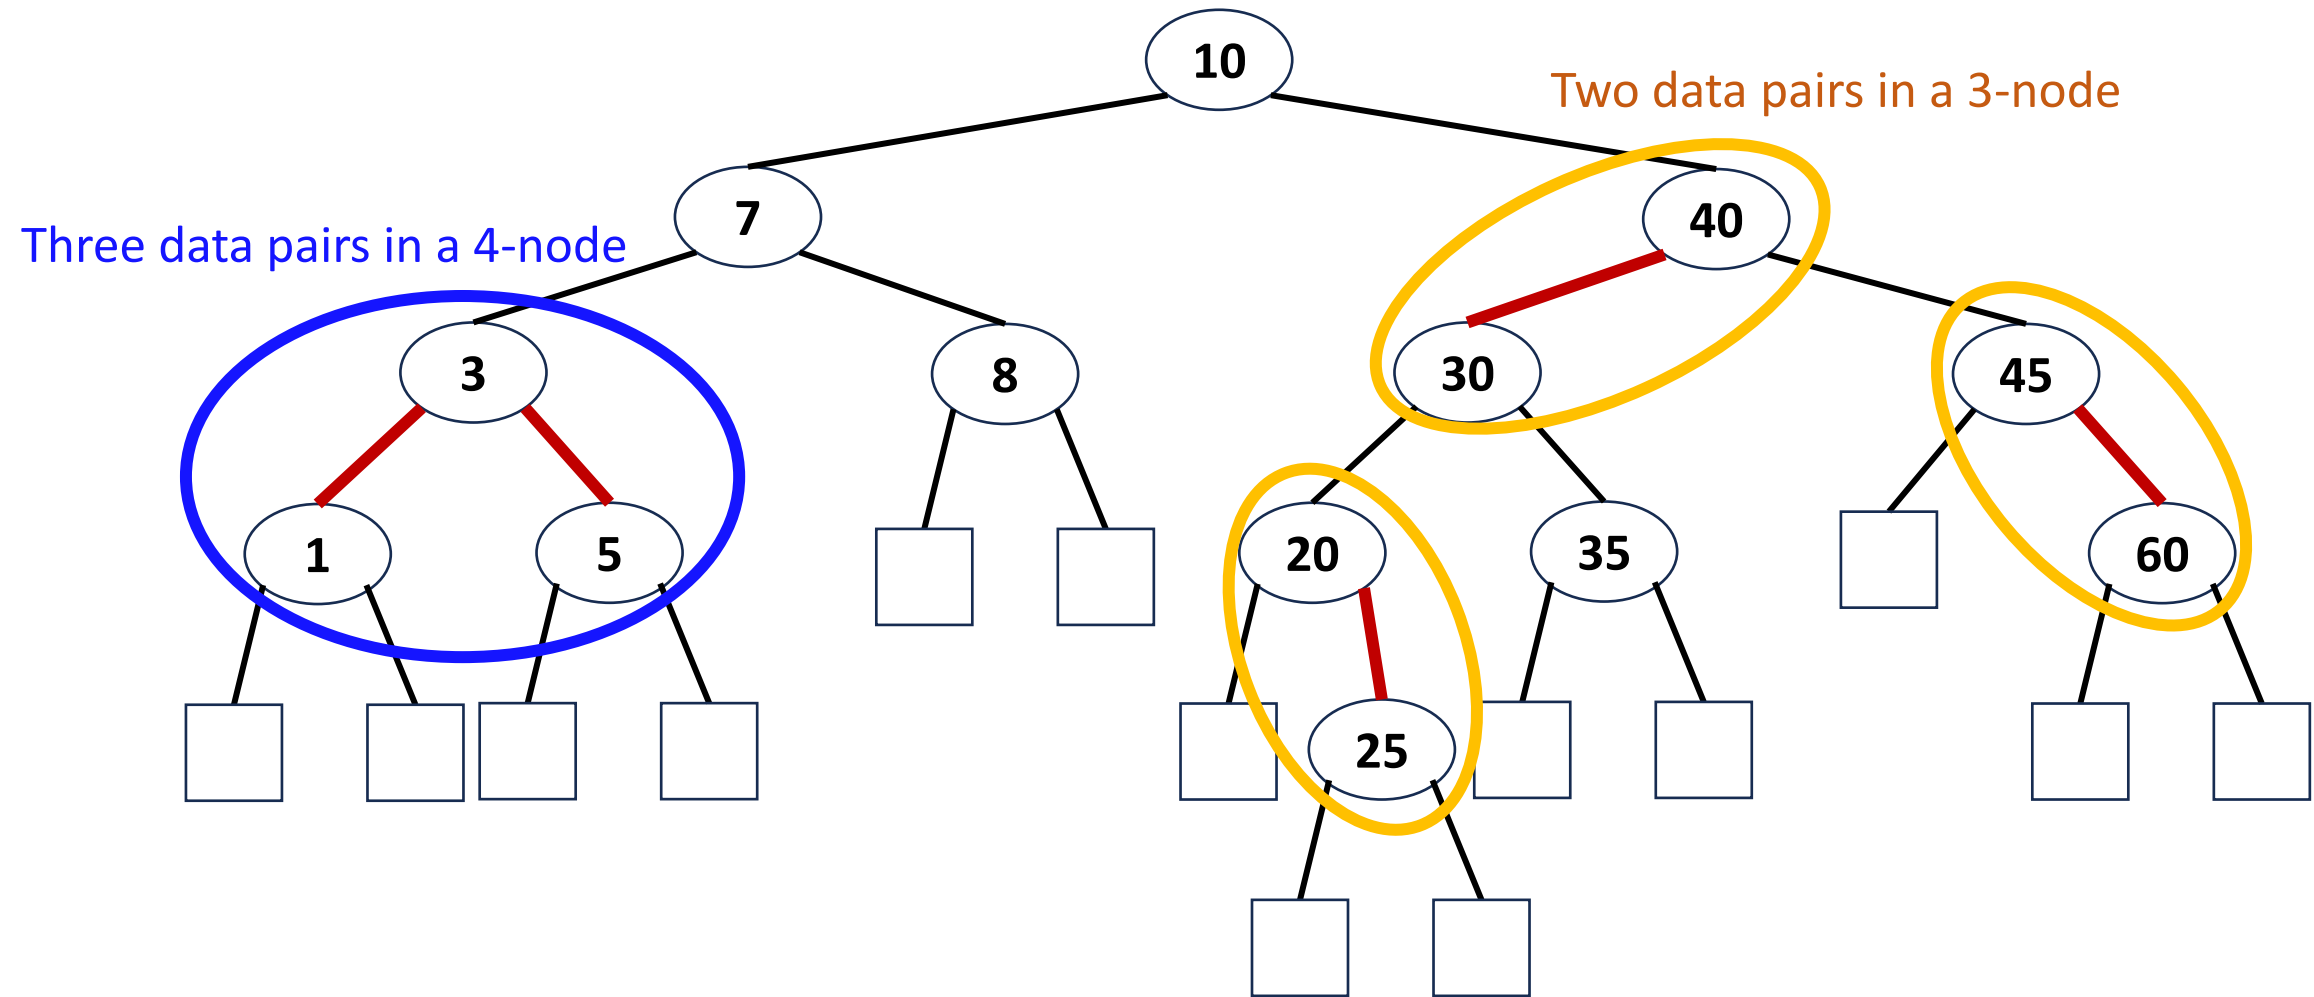


Red-black trees are just a way to represent 2-3-4 trees.

Relate 2-3-4 tree to red-black tree

- Red-black search tree uses binary trees representing m-way search trees.
- Motivation: We need a way to quickly search the data pairs within a node of m-way search trees.

Example of red-black tree



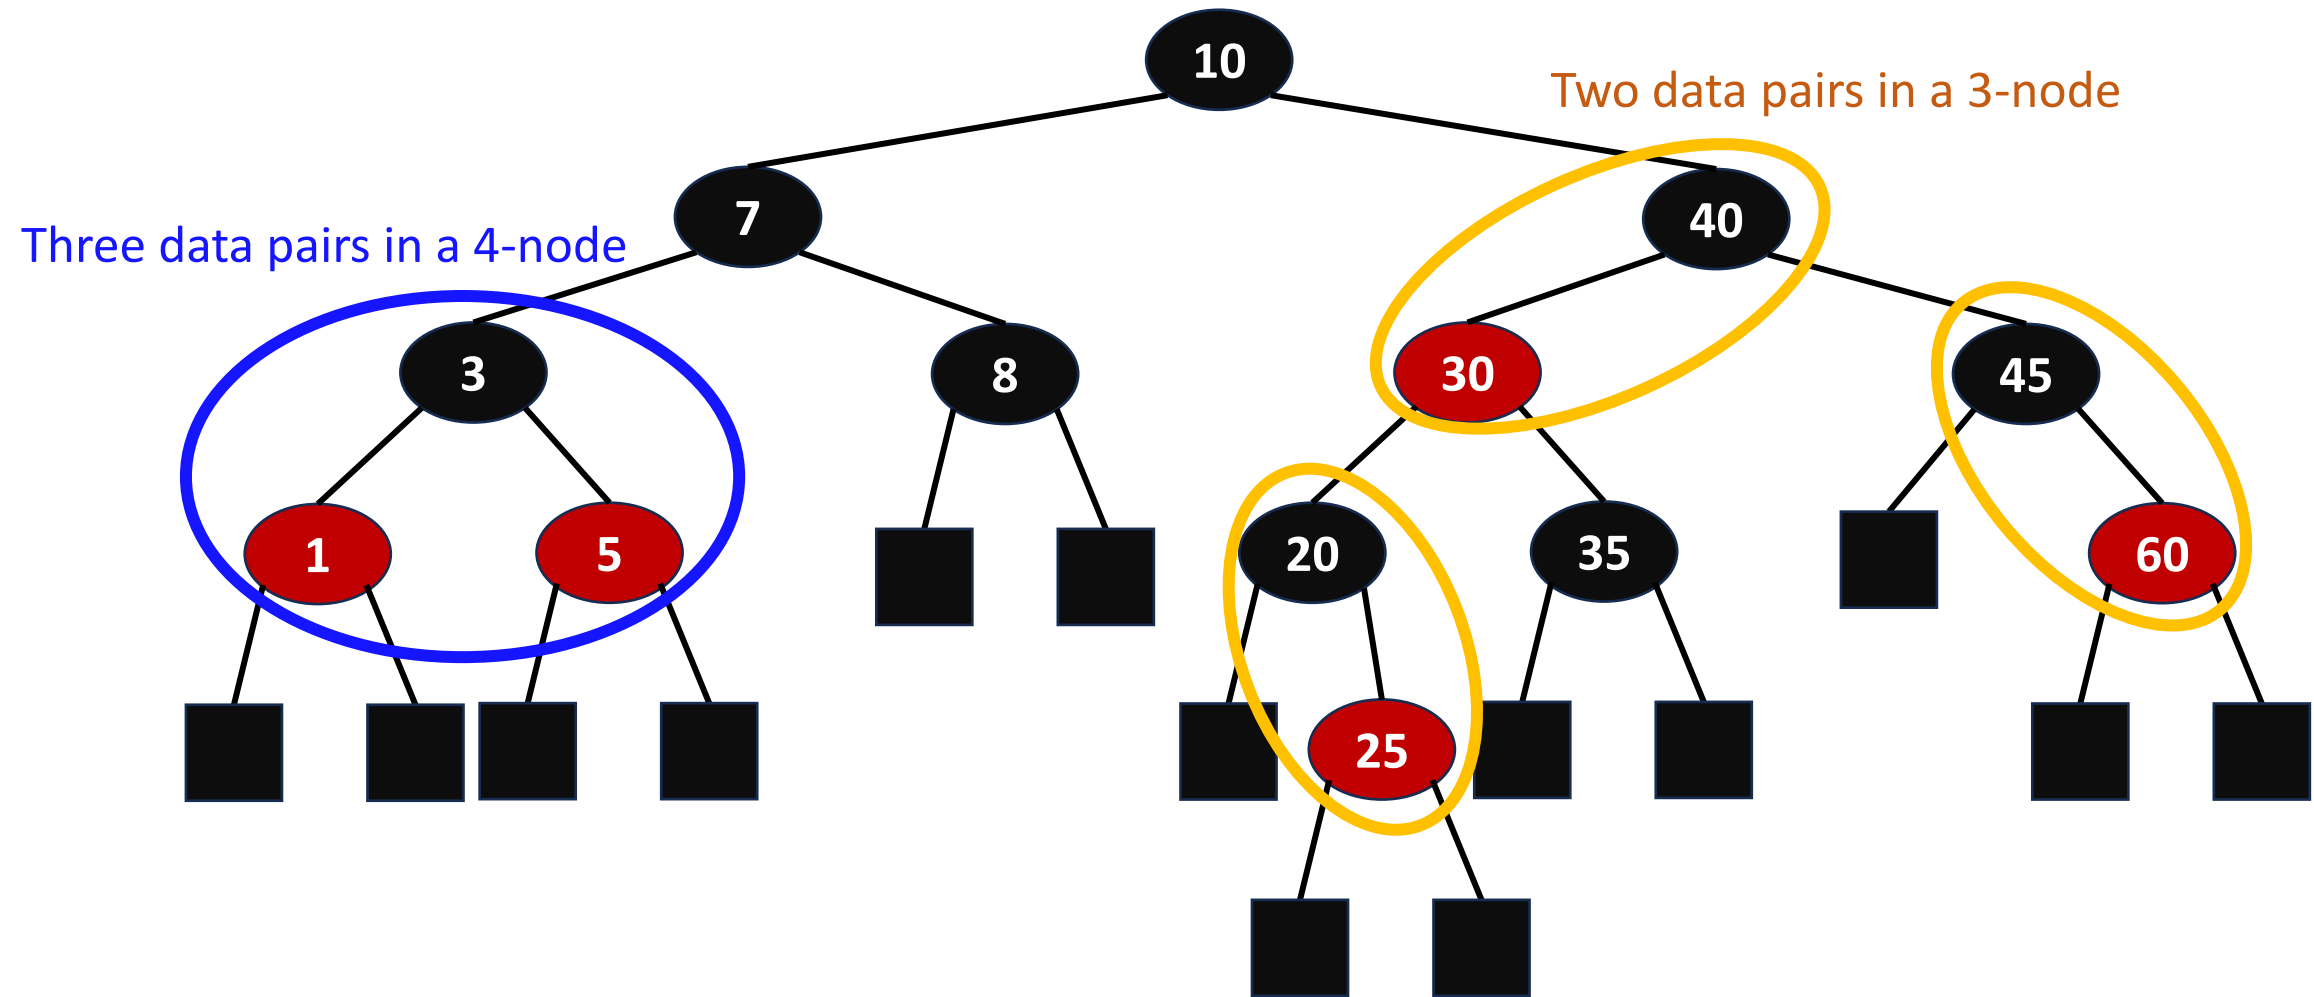
Operations for Red-black trees can be applied to 2-3-4 trees.

Red-Black tree (Equivalent definition)

A binary search tree with the following properties:

- Nodes are colored **red** or **black**.
- ***No two consecutive red nodes*** on any root-external node path
- ***Same number of black nodes*** on any root-external node path
- ***Root and external nodes are black.***



Example of red-black tree (Node)



If we know the edge colors, we can deduce the node colors and vice versa.

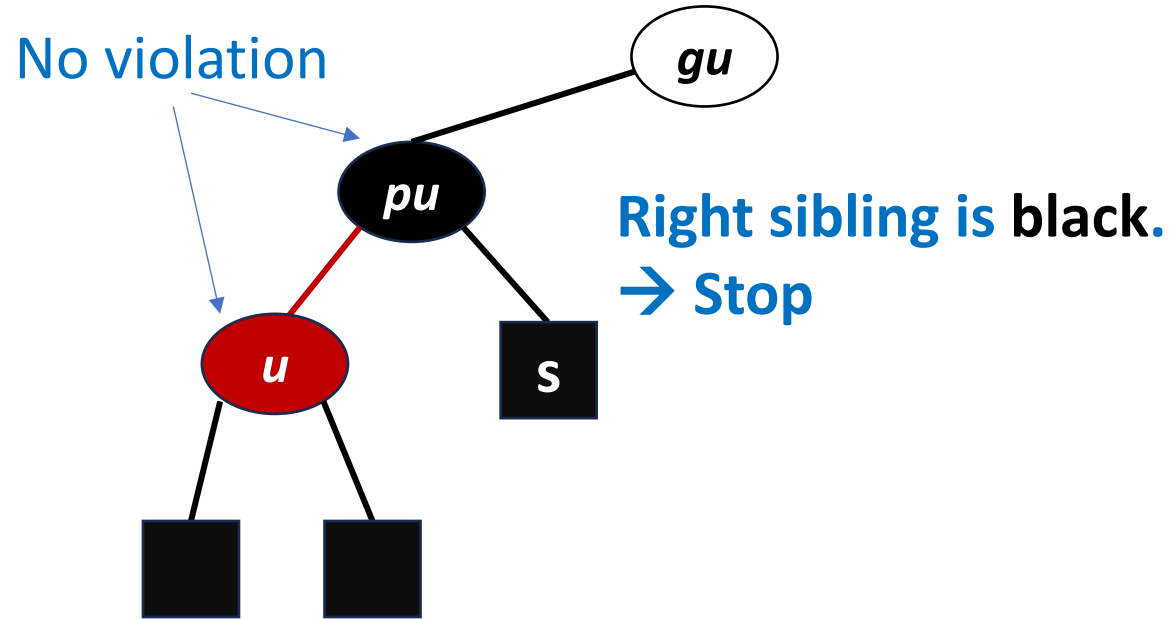
Operation: Insertion

1. Perform a standard search to find the external node where the key should be added.
2. Replace the external node with an internal node with the new key.
3. Color the new node **red**.
4. Add two new external nodes and color them **black**.
5. If the parent is **red**, we have two consecutive **red** nodes. The tree has to be reorganized.




Insertion: parent is black

- u : the new node
- pu : parent of u node
- gu : grandparent of u node
- s : sibling of u node
(may be an external node)

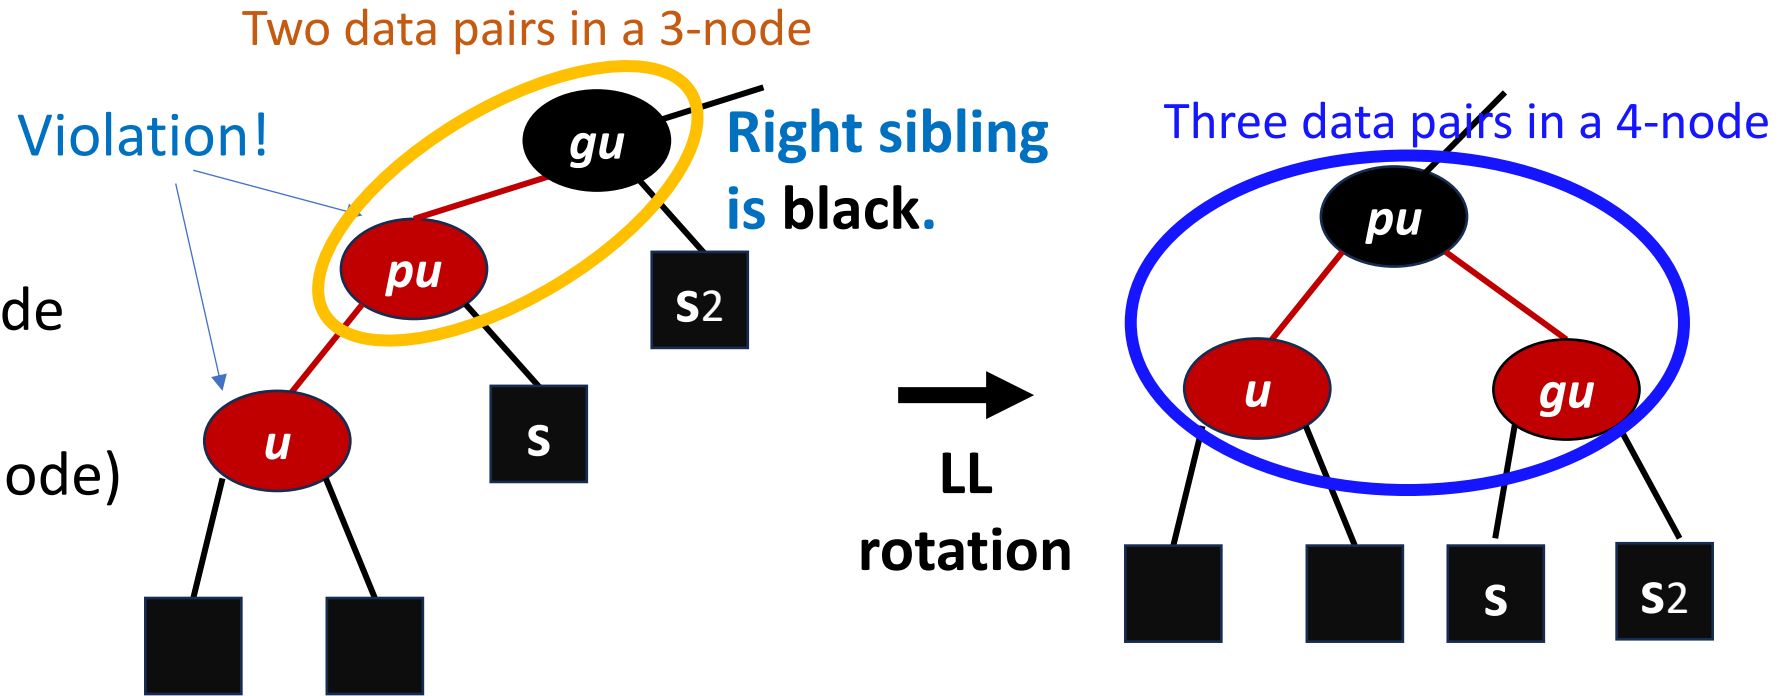


- Equivalent to insert a data pair into a 2-node of 2-3-4 tree




Insertion: parent is red and its sibling is black (1)

- u : the new node
- pu : parent of u node
- gu : grandparent of u node
- s : sibling of u node
(may be an external node)




- Equivalent to insert a data pair into a 3-node of 2-3-4 tree and become a 4-node




Insertion: parent is red and its sibling is black (2)


- u : the new node
- pu : parent of u node
- gu : grandparent of u node
- s : sibling of u node
(may be an external node)




- Equivalent to insert a data pair into a 3-node of 2-3-4 tree and become a 4-node




Insertion: parent is red and its sibling is black (3)



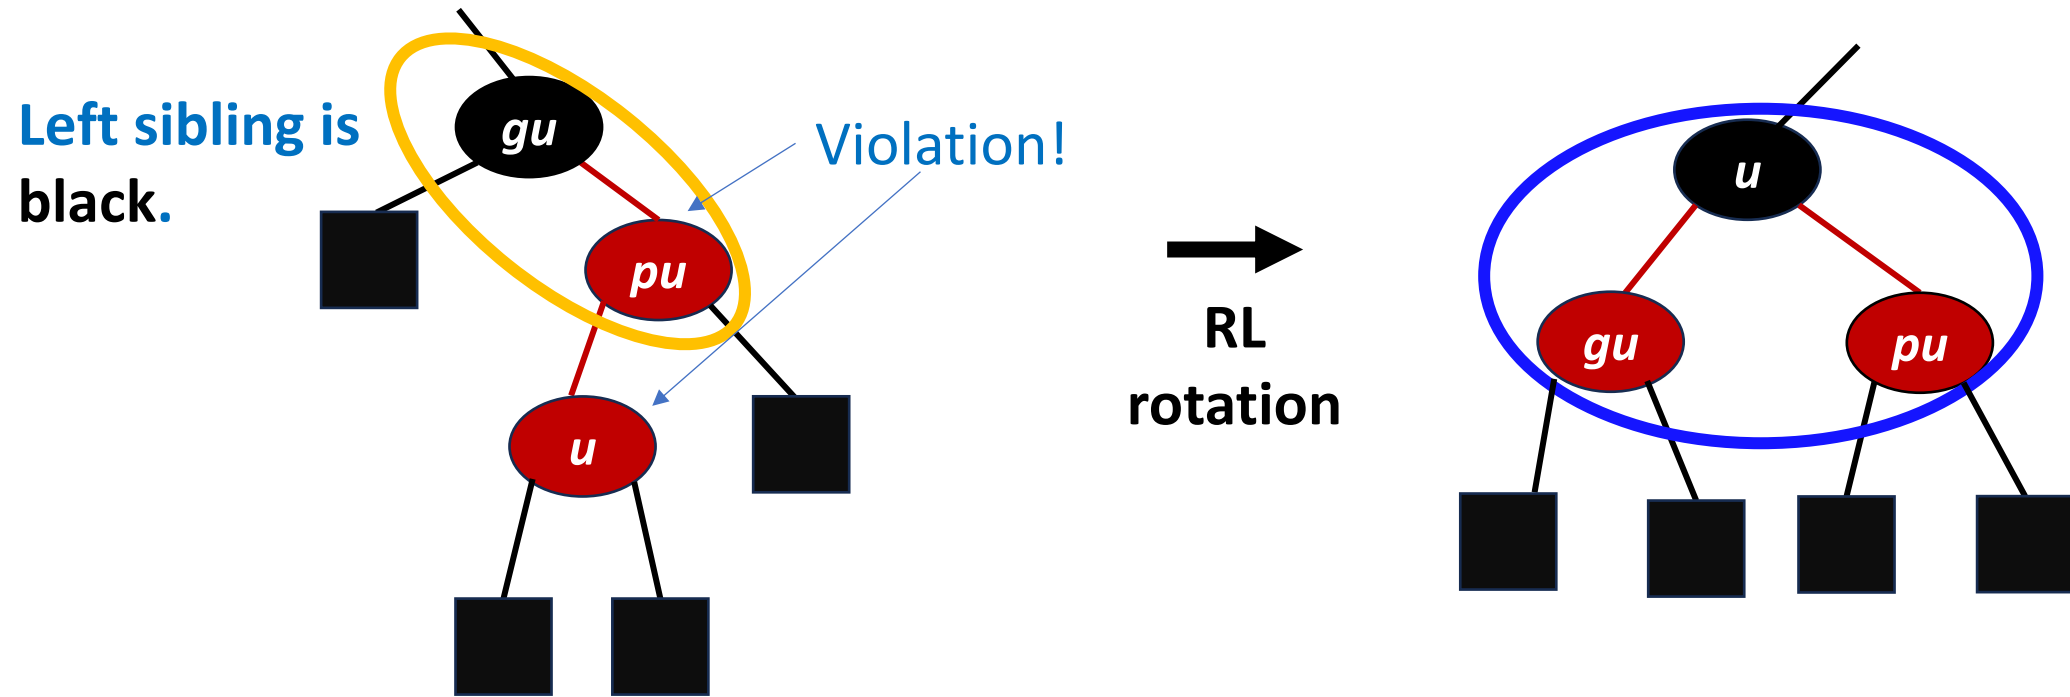
→
LR
rotation
(RR+LL)




- Equivalent to insert a data pair into a 3-node of 2-3-4 tree and become a 4-node



Insertion: parent is red and its sibling is black (4)



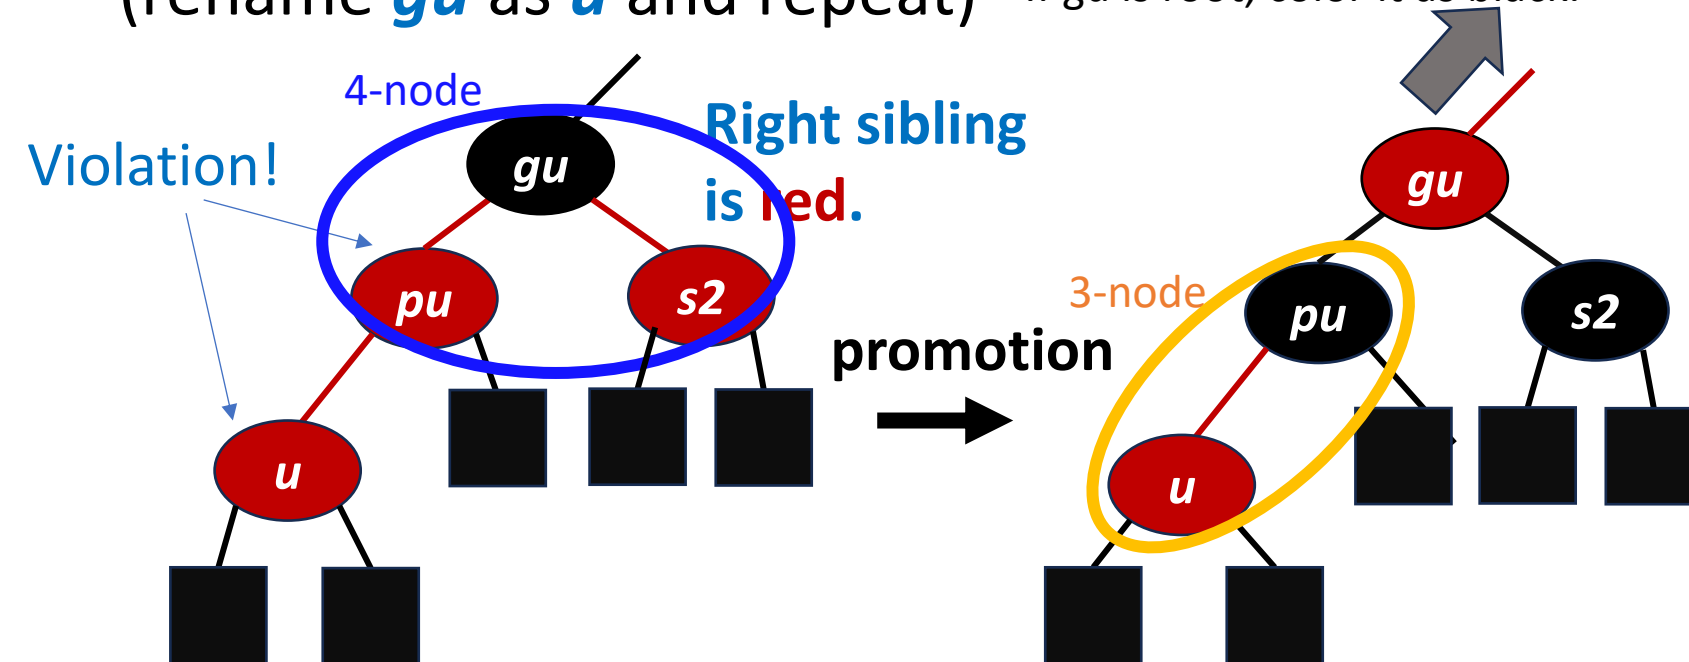
- Equivalent to insert a data pair into a 3-node of 2-3-4 tree and become a 4-node



Insertion: parent is red and its sibling is red


Bottom-up rebalancing

- Change the color of the parent *pu*, and parent's sibling *s2* to **black**.
- Change the color of the grandparent *gu* to **red**.
- Continue upward beyond *gu* if necessary.
(rename *gu* as *u* and repeat) If *gu* is root, color it as black.



Note that the *black depth* remains unchanged for all the descendants of *gu*.

Equivalent to insert a data pair into a 4-node of 2-3-4 tree. Thus, **split** is required.



Performance

- Lemma 10.1: Let P and Q be the two root-to-external node paths in a red-black tree. Then $\text{length}(P) \leq 2\text{length}(Q)$.

*Length: The number of edges on the path

Red, black, red, black,
..., red, and black

All edges are
black.

- Lemma 10.2: Let h be the height of a red-black tree, n be the number of internal nodes, and r be the rank of the root.

a) $h \leq 2r$

Lemma 10.1

b) $n \geq 2^r - 1$

If all internal nodes are black, the number of internal nodes n is $2^r - 1$.
 $r \leq \log_2(n+1)$

c) $h \leq 2\log_2(n+1)$

Combine a) and b)


*Rank: The number of black edges on the path to external nodes.

Time complexity of insertion depends on tree height h : $O(\log n)$


Example of insertion

- Start by inserting “REDSOX” into an empty tree

Red-black tree




2-3-4 tree



- Let's insert C, U, B, and S


Example of insertion

- Insert C




Parent is **black**.

2-3-4 tree



Example of insertion


- Insert U



Parent is **red** and its sibling is **black**.


→ RL rotation

2-3-4 tree




Example of insertion

- Insert B



2-3-4 tree

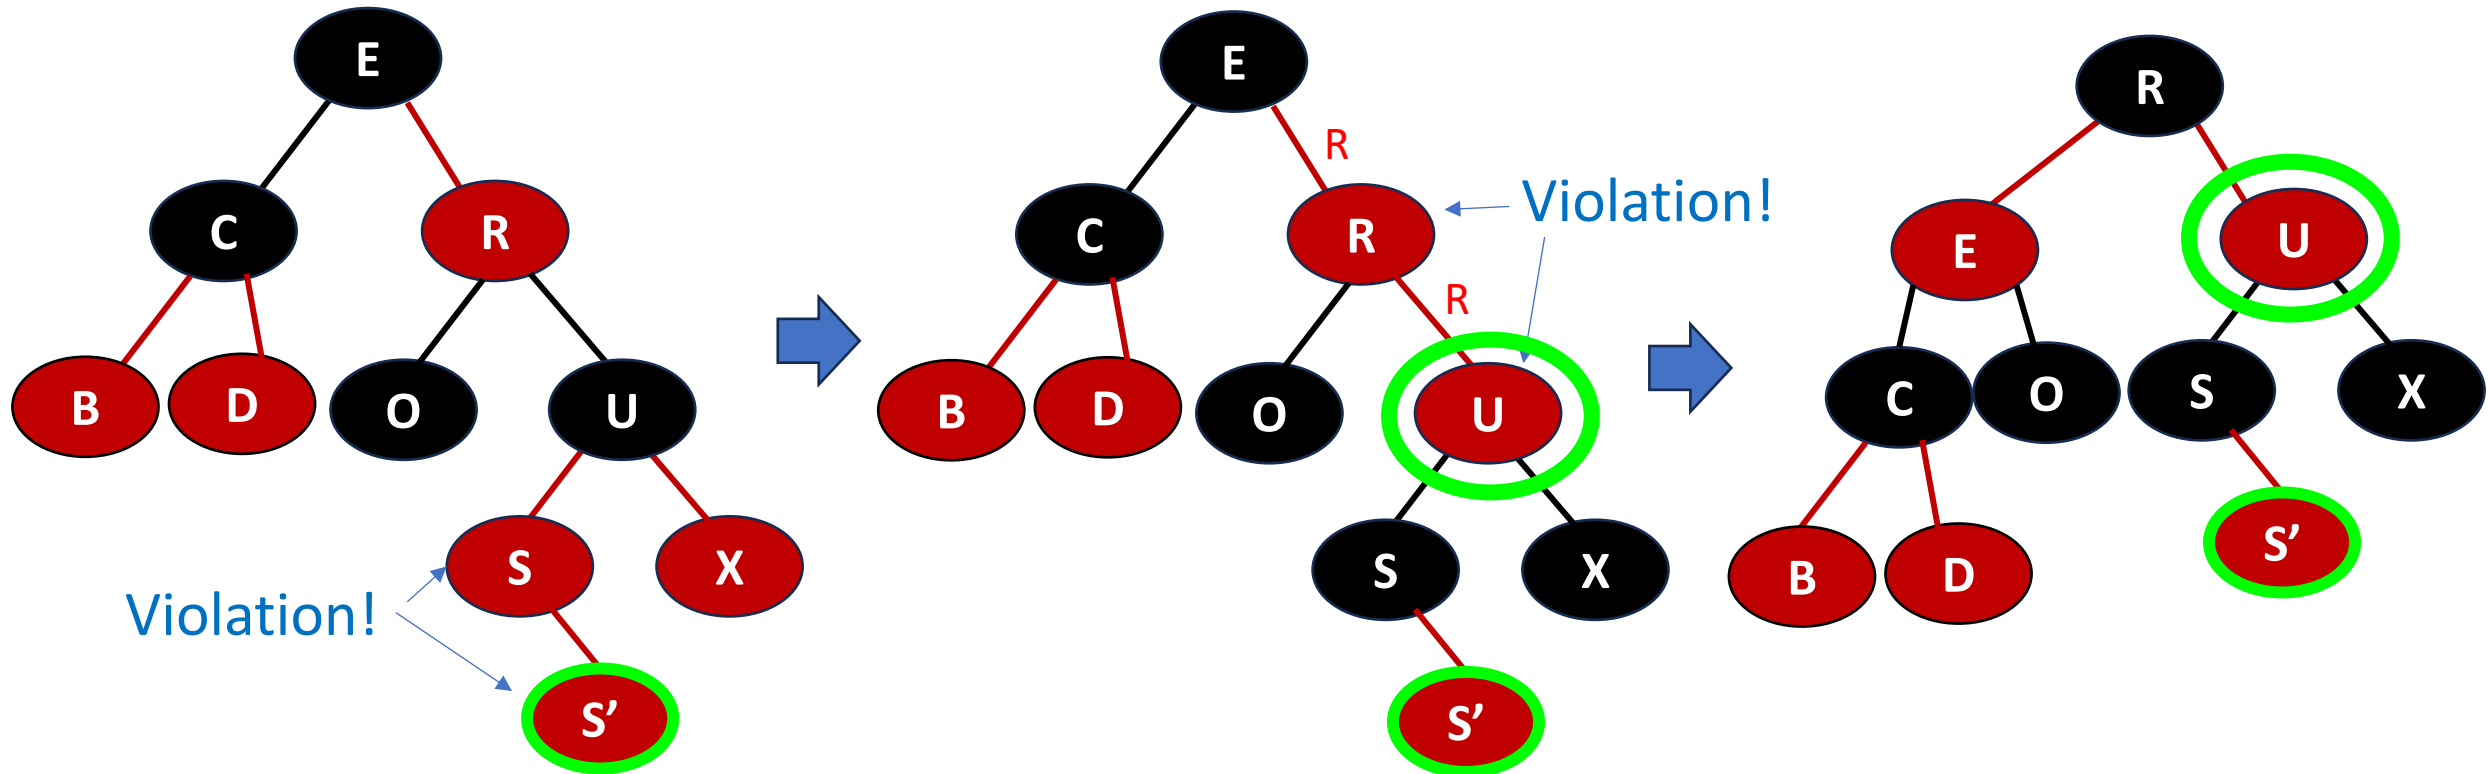


Parent is **red** and its sibling is **black**.

→ LL rotation

Example of insertion

- Insert S



It can be on the left or right.

Parent is **red** and its sibling is **red**. Parent is **red** and its sibling is **black**.


→ promotion (color change)

→ RR rotation

Example of insertion


- Insert S

2-3-4 tree



Exercise

- Given the following red-black tree.
 - Q1: Please insert 70. What is(are) the key(s) in red nodes?
 - Q2: (Continue 1) Please insert 60. What is(are) the key(s) in red nodes?
 - Q3: (Continue 2) Please insert 65. What is(are) the key(s) in red nodes?
 - Q4: (Continue 3) Please insert 62. What is(are) the key(s) in red nodes?

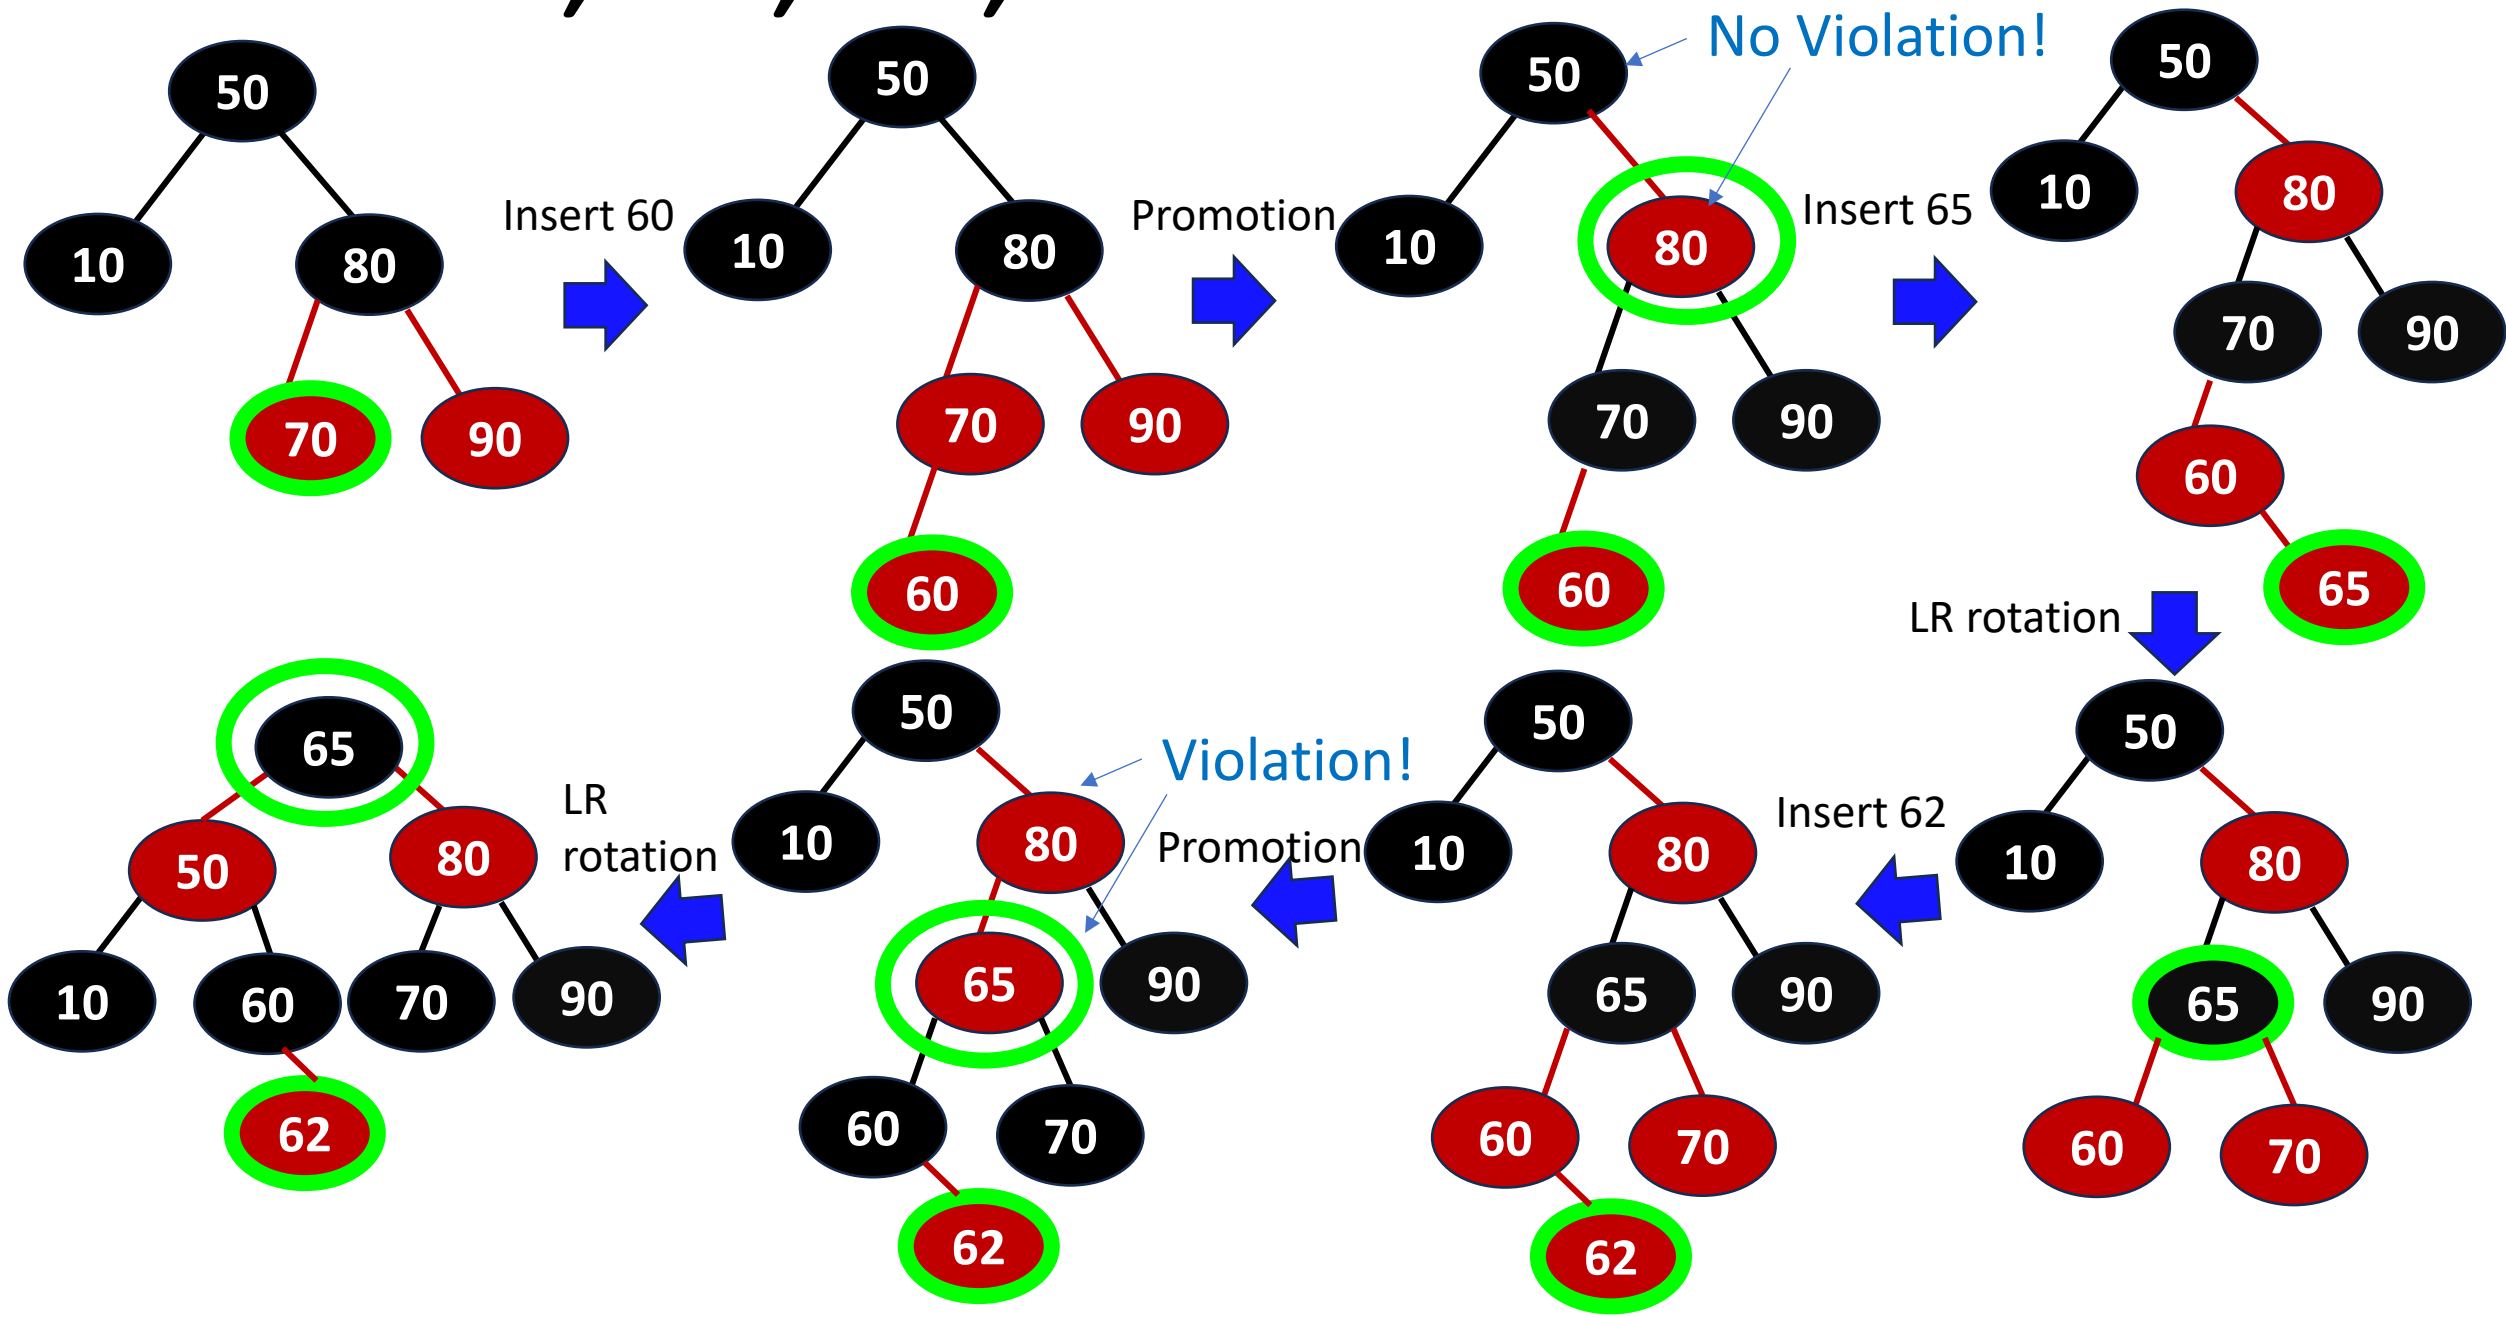


Please reply your answers of Q1-Q4 via the following link:




<https://forms.gle/erneNgrctngmjem1A>
Group members: 2~4 people

Insert 70, 60, 65, 62




Operation: Deletion

- An exercise in the textbook
- We will only discuss about the deletion of nodes nearby leaf nodes.
 - Reason: The node to replace the deleted node is always from a leaf node in 2-3-4 tree.



The inorder predecessor always has at most one child, since it has the largest key in the left subtree.
The inorder successor always has at most one child, since it has the smallest key in the right subtree.


Inorder predecessor of 11




Inorder successor of 11

Another case: when 7 is red
After transformation, 7 will become in a black node and we don't have to change the structure.

Delete 11

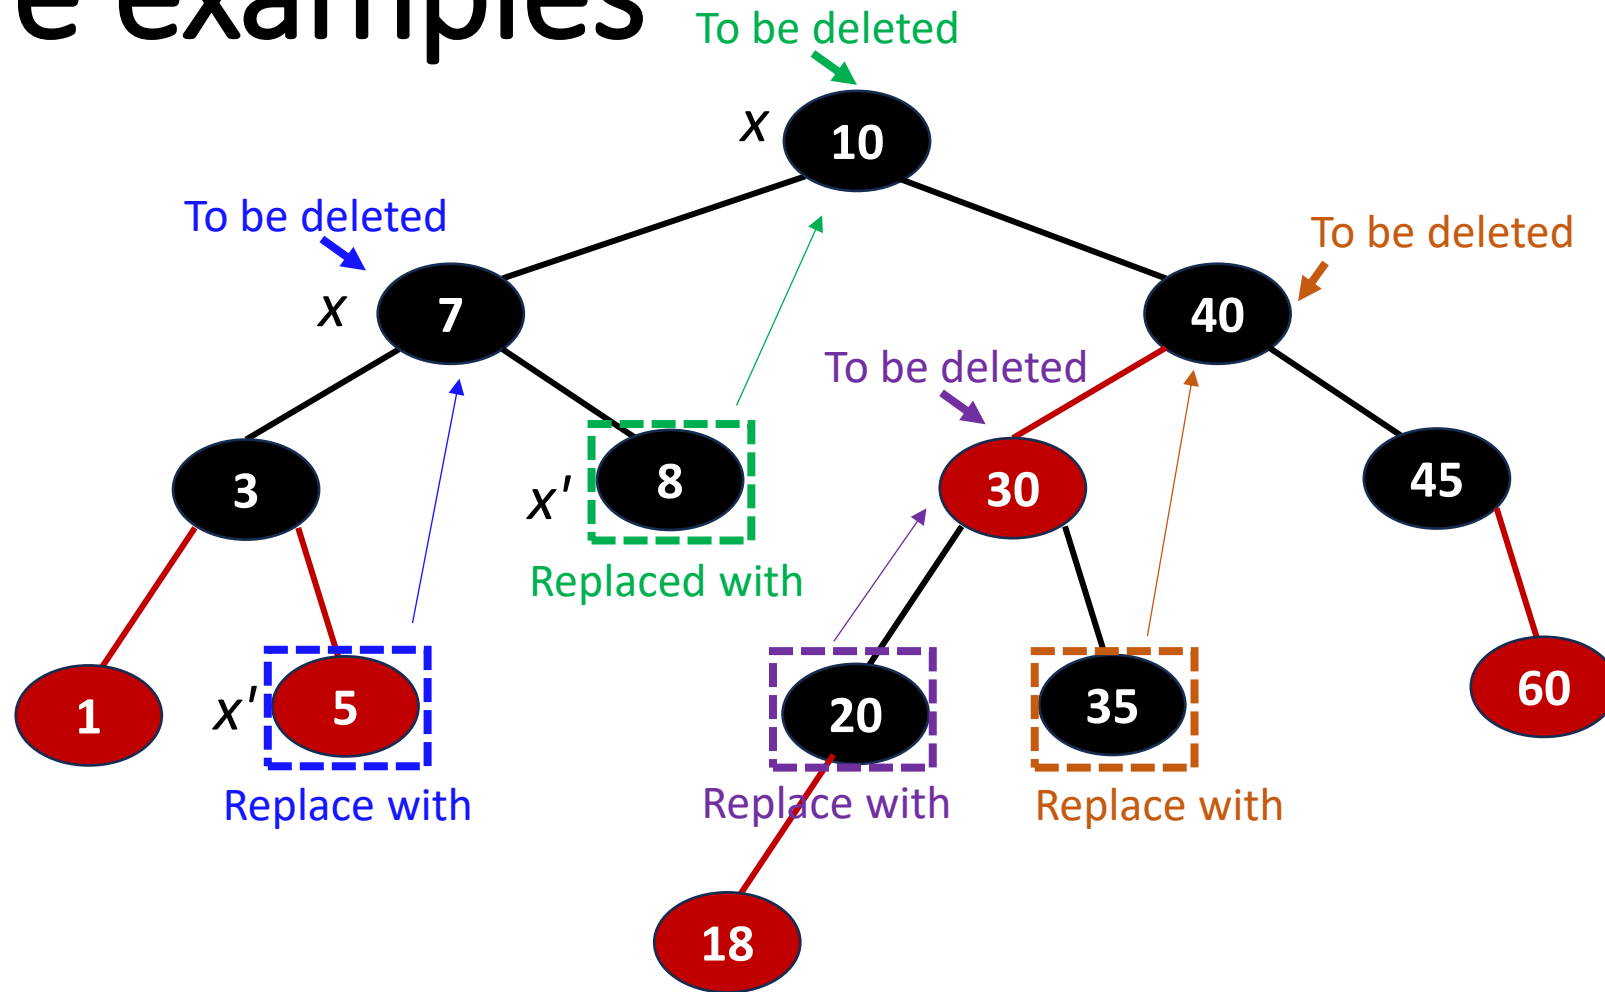


Transform to



Become the deletion in case y1
(it will be discussed later).

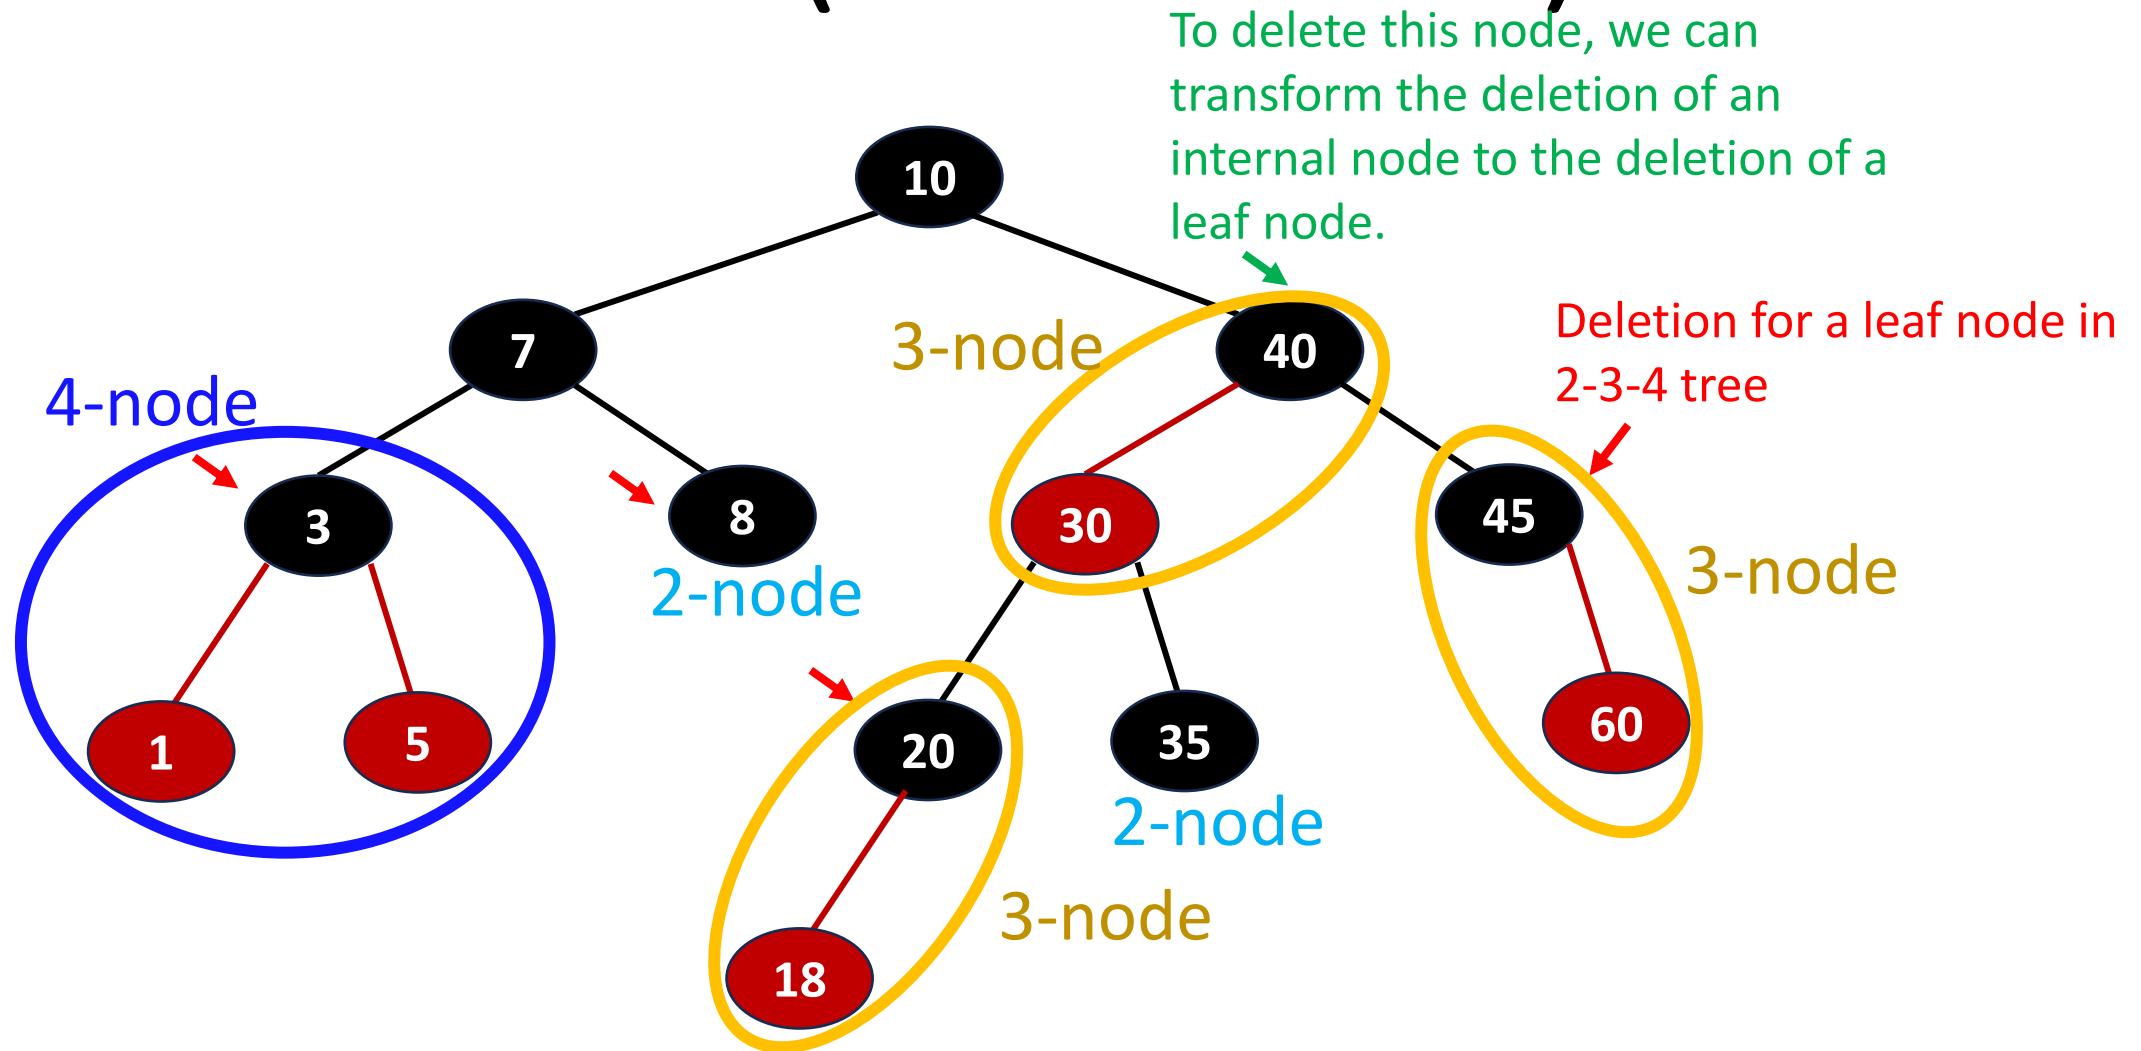
More examples



Delete x and replace with x' :


After transformation, x' becomes the node to be deleted and it is always in the leaf node of a 2-3-4 tree.

Cases of deletion (in 2-3-4 tree)



Recall: How do we delete the data pair in the 2-node, 3-node, and 4-node in 2-3-4 tree?


Delete y




y_0 , y_1 , y_2 are different types of nodes to be deleted:

- y_0 : the node connects its parent with a **red link** or it's a red node.
 - It's within a 3- or 4-node in 2-3-4 tree. → Simply delete y .
- y_1 and y_2 : the nodes connect to its parent with a **black link**.
 - y_1 : with red child node. It's within a 3- or 4-node in 2-3-4 tree.
 - y_2 : without red child node. It affects balance in 2-3-4 trees and needs reshape.


Delete y_0



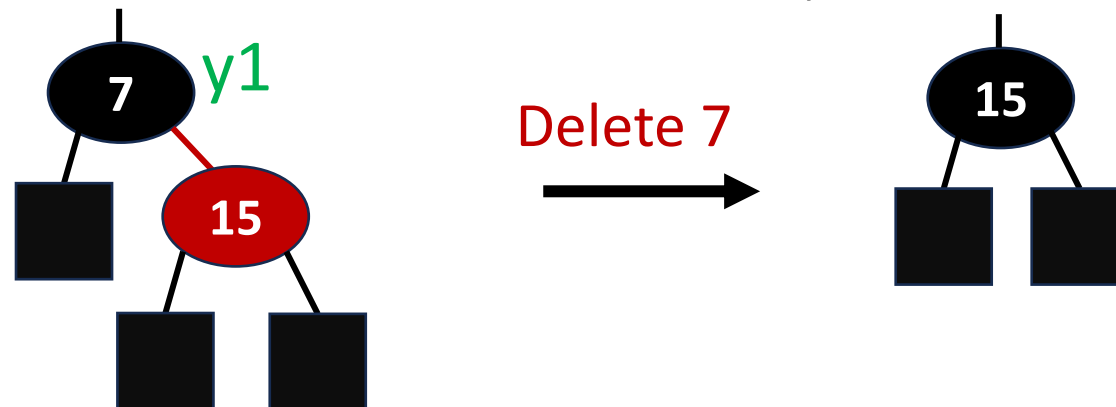
- y_0 : the node connects its parent with a red link or it's a red node.
- Delete y_0 and promote one of its external nodes.



Delete **y1**




- **y1**: it's a black node with at least one red child node. Within a 3- or 4-
node in 2-3-4 tree.
- Delete **y1**, promote one of its child nodes, which is then colored black.





Example of **y1** deletion

Delete 11




Transform to



Cases of y_2

y_2 : it's a black node without red child node.




- Delete y :
 - It's better that the tree height remains the same.
 - If the reshaping propagates towards root, tree height may be reduced by one.
- Before reshaping, it's a b-tree (2-3-4 tree). After reshaping, it's still a b-tree.

Type I of y_2

Recall: In m-way search tree, **rotation** has high priority. Otherwise, **combine** may reduce the height by 1.


- Sibling of y is within a 3- or 4- node in a 2-3-4 tree
 - It can provide one data pair.
 - Rotation operation



Type II of y_2

Recall: In m-way search tree, **rotation** has high priority. Otherwise, **combine** may reduce the height by 1.



- Sibling of y is within a 2-node in 2-3-4 tree
 - It cannot provide one data pair.
 - We cannot do rotation.
 - Combine operation



Type III of y_2

Recall: In m-way search tree, **rotation** has high priority. Otherwise, **combine** may reduce the height by 1.

- In red-black tree, y has a red sibling s .
In 2-3-4 tree, s is not y 's sibling. Actually, y 's siblings are a and b .
→ Reshape the red-black tree, so that y 's sibling is a black node.




- The deletion may cause chain reaction.

Note of deletion

- We only adjust the nodes which are located at nearest two levels in the corresponding 2-3-4 tree.
- The adjustment may cause **chain reaction (bottom-up towards root)**.
→ Resulting from “**double black** in y”.

- Originally, a path has two consecutive black nodes.
- Because of the deletion, one black node y is removed, and the other node is moved to the location of y.
- The node y is marked as **double black** to show that it should have two black nodes to maintain the black height.

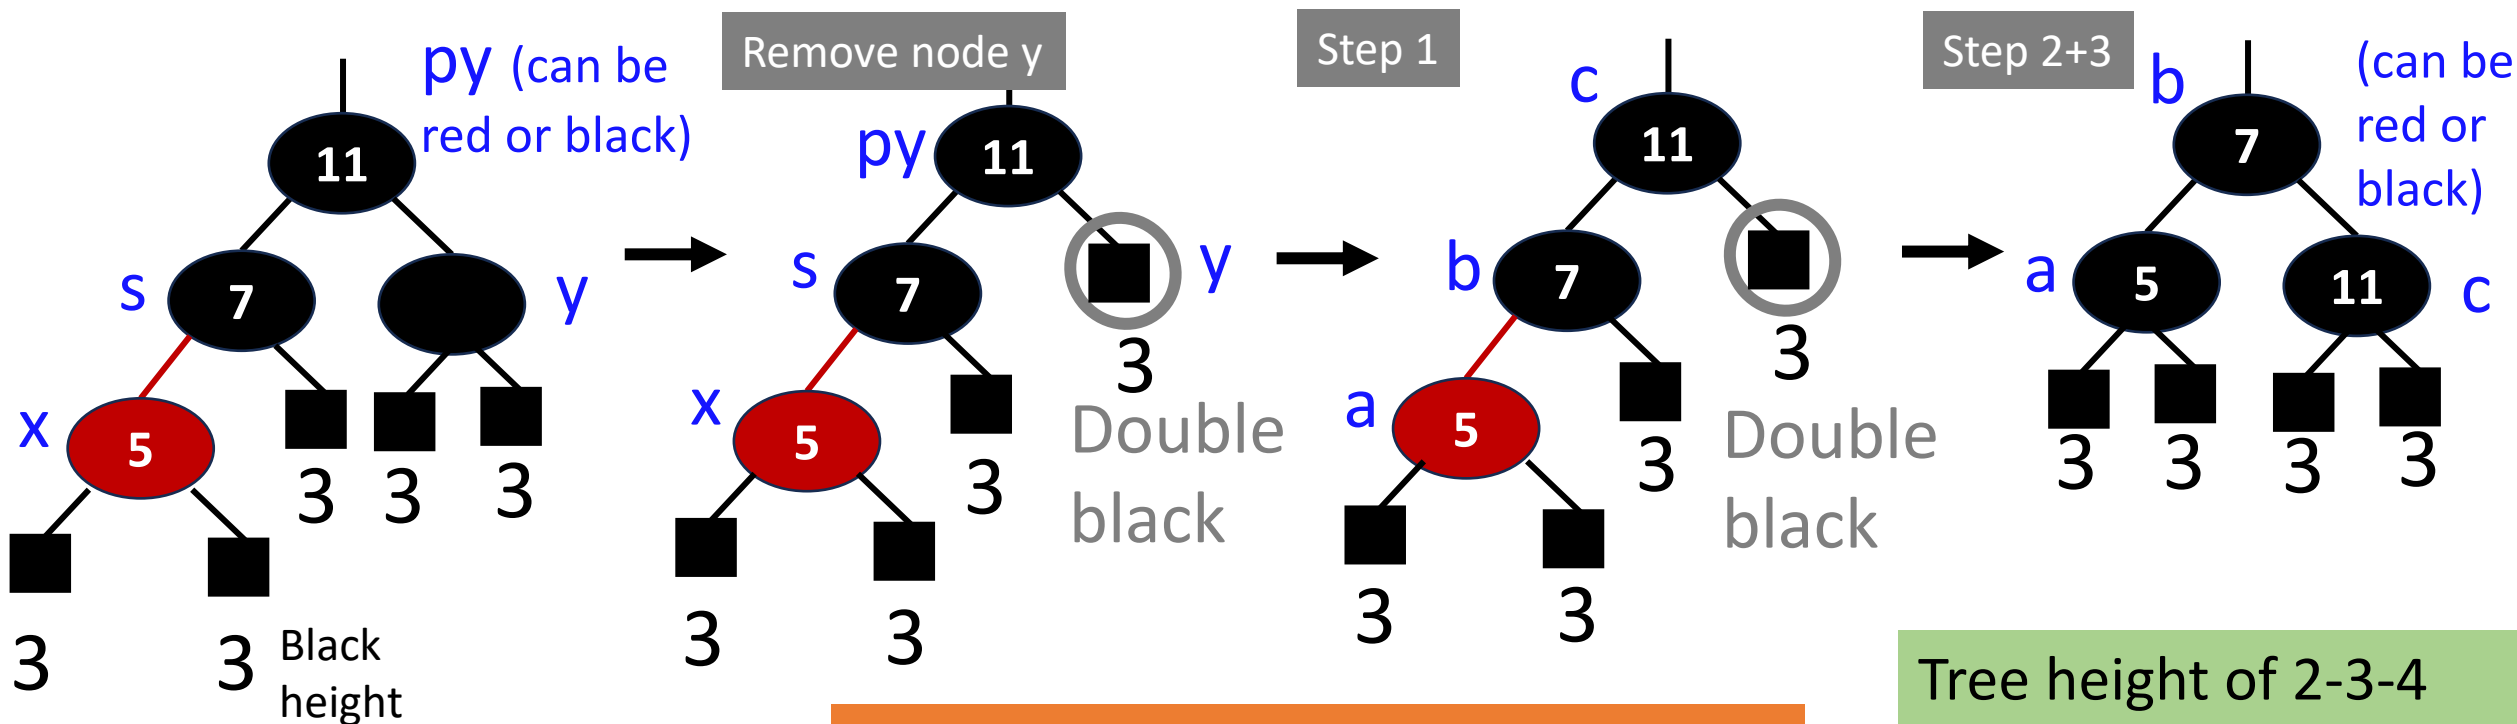


Type I of y_2

Type I: Sibling s is **black** and has a **red child** x .

Generally, relabel to let key $a < \text{key } b < \text{key } c$

1. Relabel nodes: $x \rightarrow a$, $s \rightarrow b$, and $py \rightarrow c$ (For LL case in this example)
2. Replace the original py with b . Make a and c its children. Keep inorder relationships unchanged.
3. Color a and c **black**, and color b the former color of py .



y 's sibling can provide a data pair.


Tree height of 2-3-4 tree is not changed.

Type II of y_2


Type II: Sibling s is **black** and has **black** children.

1. Color s **red**.
2. If py is **red**, we color py **black**.
3. Else if py is not the root, we color py **double black** and repeat resolving for py .

May repeat Type I,
Type II, Type III of y_2



Tree height of 2-3-4 tree is not changed.




Tree height of 2-3-4 tree reduces by one if py is root.

May result in chain reaction.


y 's sibling cannot provide a data pair. Need the data pair from y 's ancestor.

Type II of y_2


Type II: Sibling s is **black** and has **black** children.




Tree height of 2-3-4 tree is not changed.



In 2-3-4 tree, removing y represents removing a child from y 's parent. We borrow a data pair from the parent node and combine with the sibling.



If a root becomes double black, the tree height of 2-3-4 tree reduces by one.



In 2-3-4 tree, after combining the sibling and the parent node, the original parent node becomes empty.

Type III of y_2

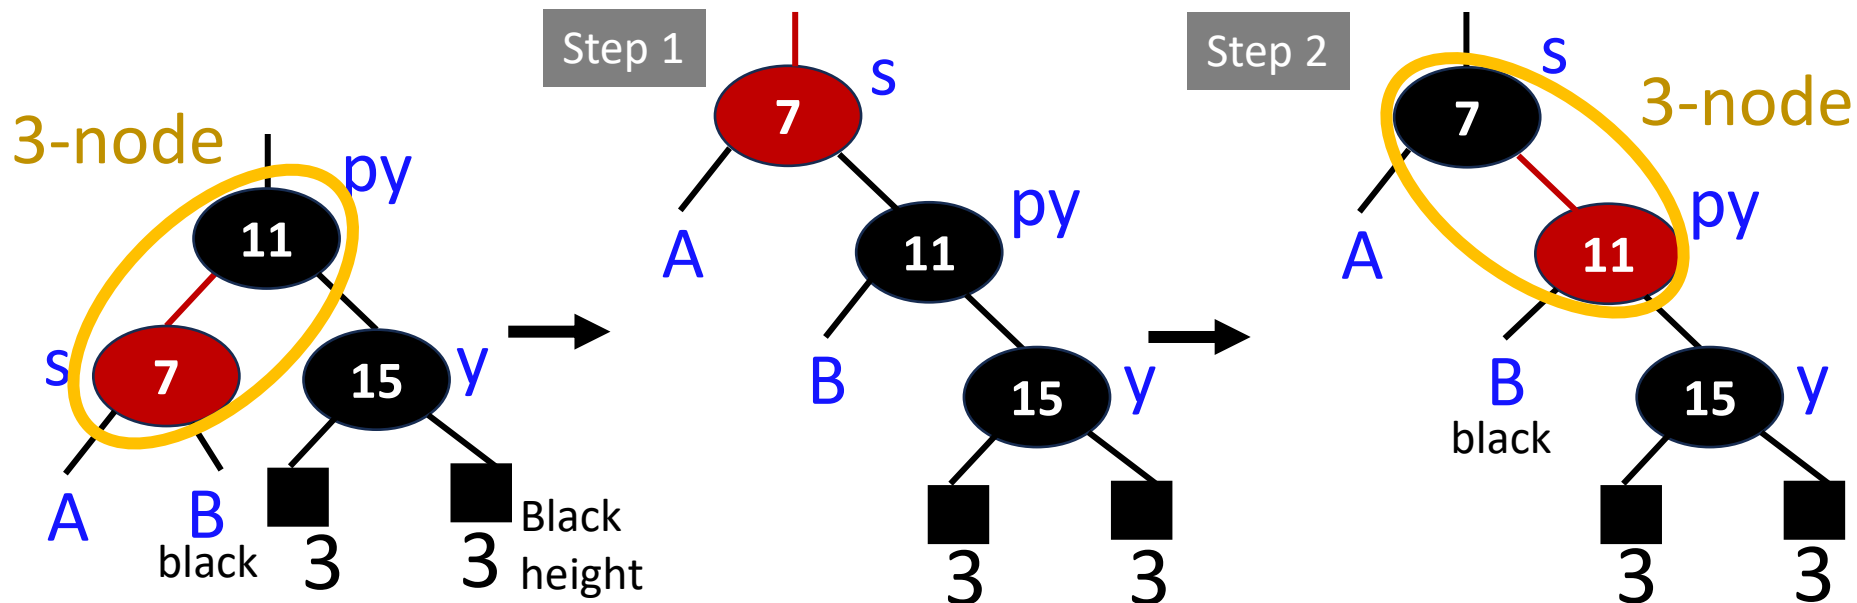

Type III: Sibling **s** is **red**.

As **s** is **red**, **py** must be **black**.

1. Rotate around **s** and **py**. That is, **s** becomes the parent of **py**.
2. Color **s** **black** and **py** **red**, and repeat resolving for **y**.

May repeat Type I and Type II of y_2 (not Type III)


Removing **y** results in double black.



In red-black tree, y has a **red** sibling. Thus, we create a **black** sibling for y based on 2-3-4 tree. Then, type III of y_2 is transformed to type I or II of y_2 .

Tree height of 2-3-4 tree may finally reduce one.

Example of deletion




Delete 5

- A **red** leaf node

→ Case **y0**

→ Simply delete 5

Example of deletion




Delete 10 (1/3)

- An internal node

→ Replace with the key of inorder predecessor, that is, 7.


Example of deletion



Delete 10 (2/3)

- An internal node
 - Replace with the key of inorder predecessor, that is, 7.
 - Delete the inorder predecessor *y*.
 - Case *y*2 (black node without a red child)
 - Type I (Sibling is black and has a red child)

Example of deletion



Delete 10 (3/3)

- An internal node

- Replace with the key of inorder predecessor, that is, 7.


- Delete the inorder predecessor y .

- Case y_2 (black node without a red child)

- Type I (Sibling is black and has a red child)

- Rotation

Example of deletion




Delete 31 (1/2)

Case y2 (black node without a red child)

Type II (Sibling is **black** and has **black** children)

→ Color its sibling red and its parent black.

Example of deletion




Delete 31 (2/2)

Case y2 (black node without a red child)

Type II (Sibling is **black** and has **black** children)

→ Color its sibling **red** and its parent **black**.

Example of deletion




Delete 7

- A **red** leaf node
- Case **y0**
- Simply delete 7


Delete 29

- A **black** leaf node **without a red child**.
- Case **y2**.
- Sibling is **red**.
- Type III of **y2**


Type III of y2
(red sibling)




1. Rotate s and py



2. Color s to black
Color py to red




Type II of y2
(Black sibling without a red child)



Combine operation: color parent py' black and sibling s' red
(Originally, py' is red. No further adjustment)

Exercise

- Given the following red-black tree.
 - Q5: Please delete 4. What is(are) the key(s) in red nodes?
 - Q6: (Continue Q5) Please delete 9. What is(are) the key(s) in red nodes?
 - Q7: (Continue Q6) Please delete 40. What is(are) the key(s) in red nodes?
 - Q8: (Continue Q7) Please delete 7. What is(are) the key(s) in red nodes?



Please reply your answers of
Q5-Q8 via the following link:




<https://forms.gle/erneNgrctngmjem1A>
Group members: 2~4 people

Exercise


- Q5: Please delete 4.

Case y0



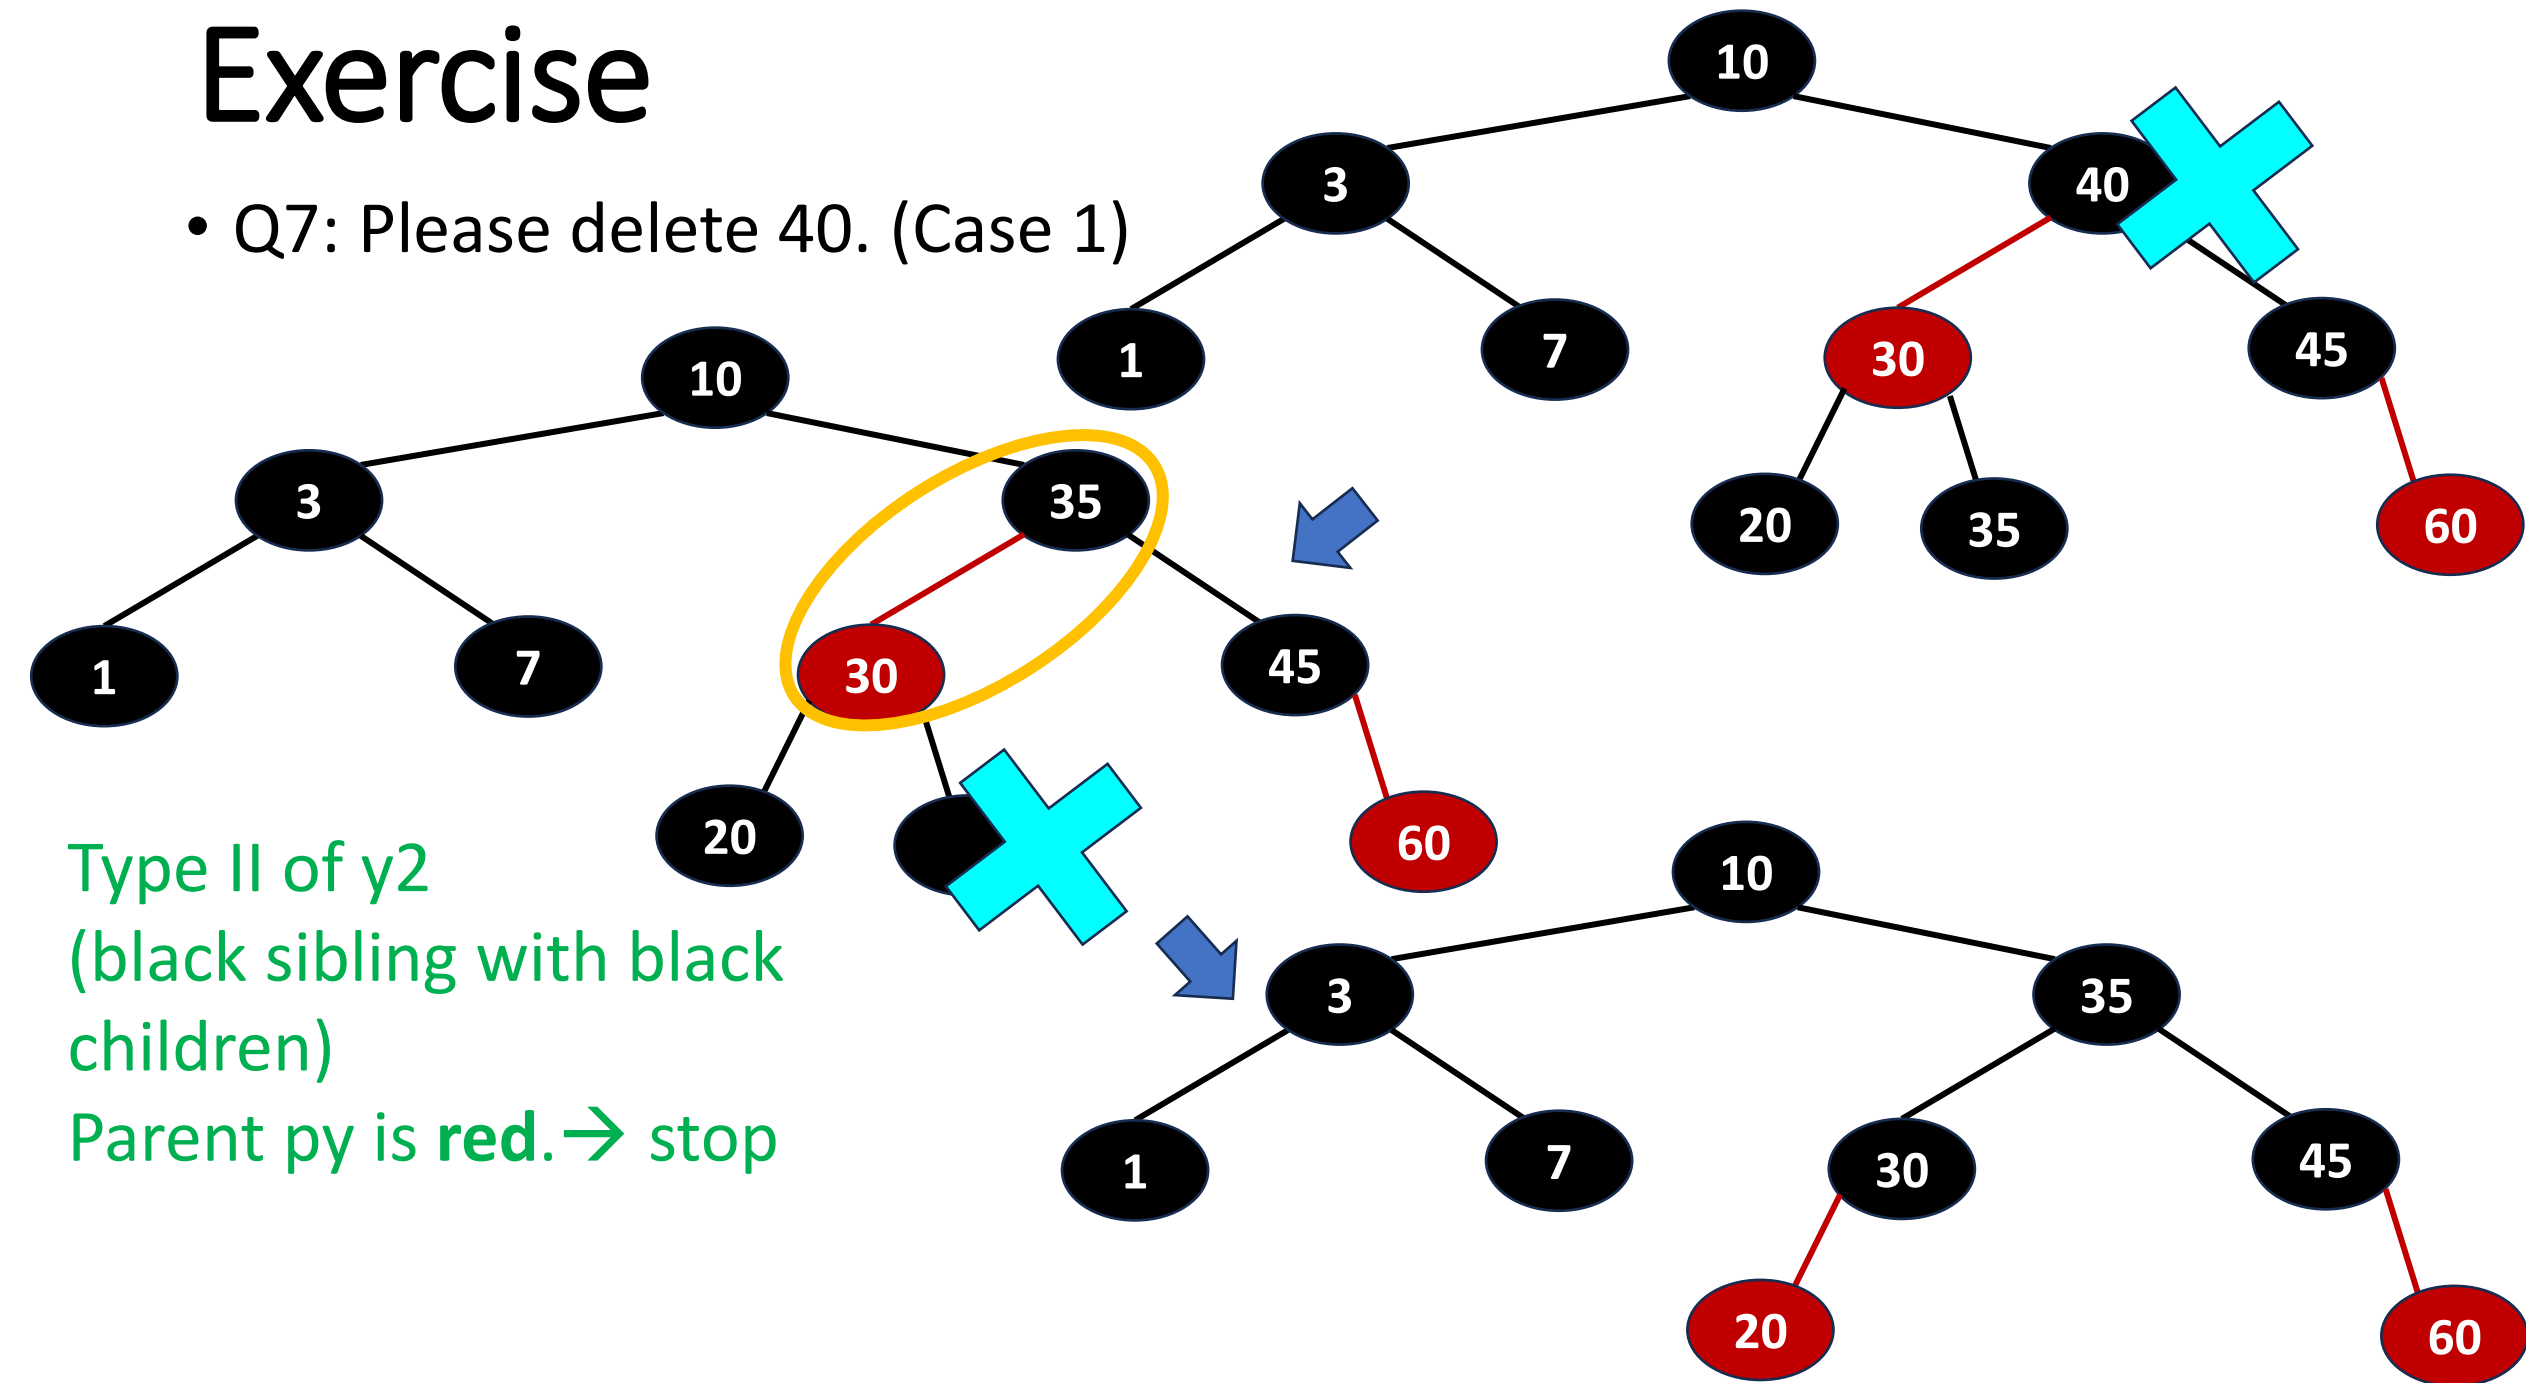
- Q6: Please delete 9.

Case y2
Type I



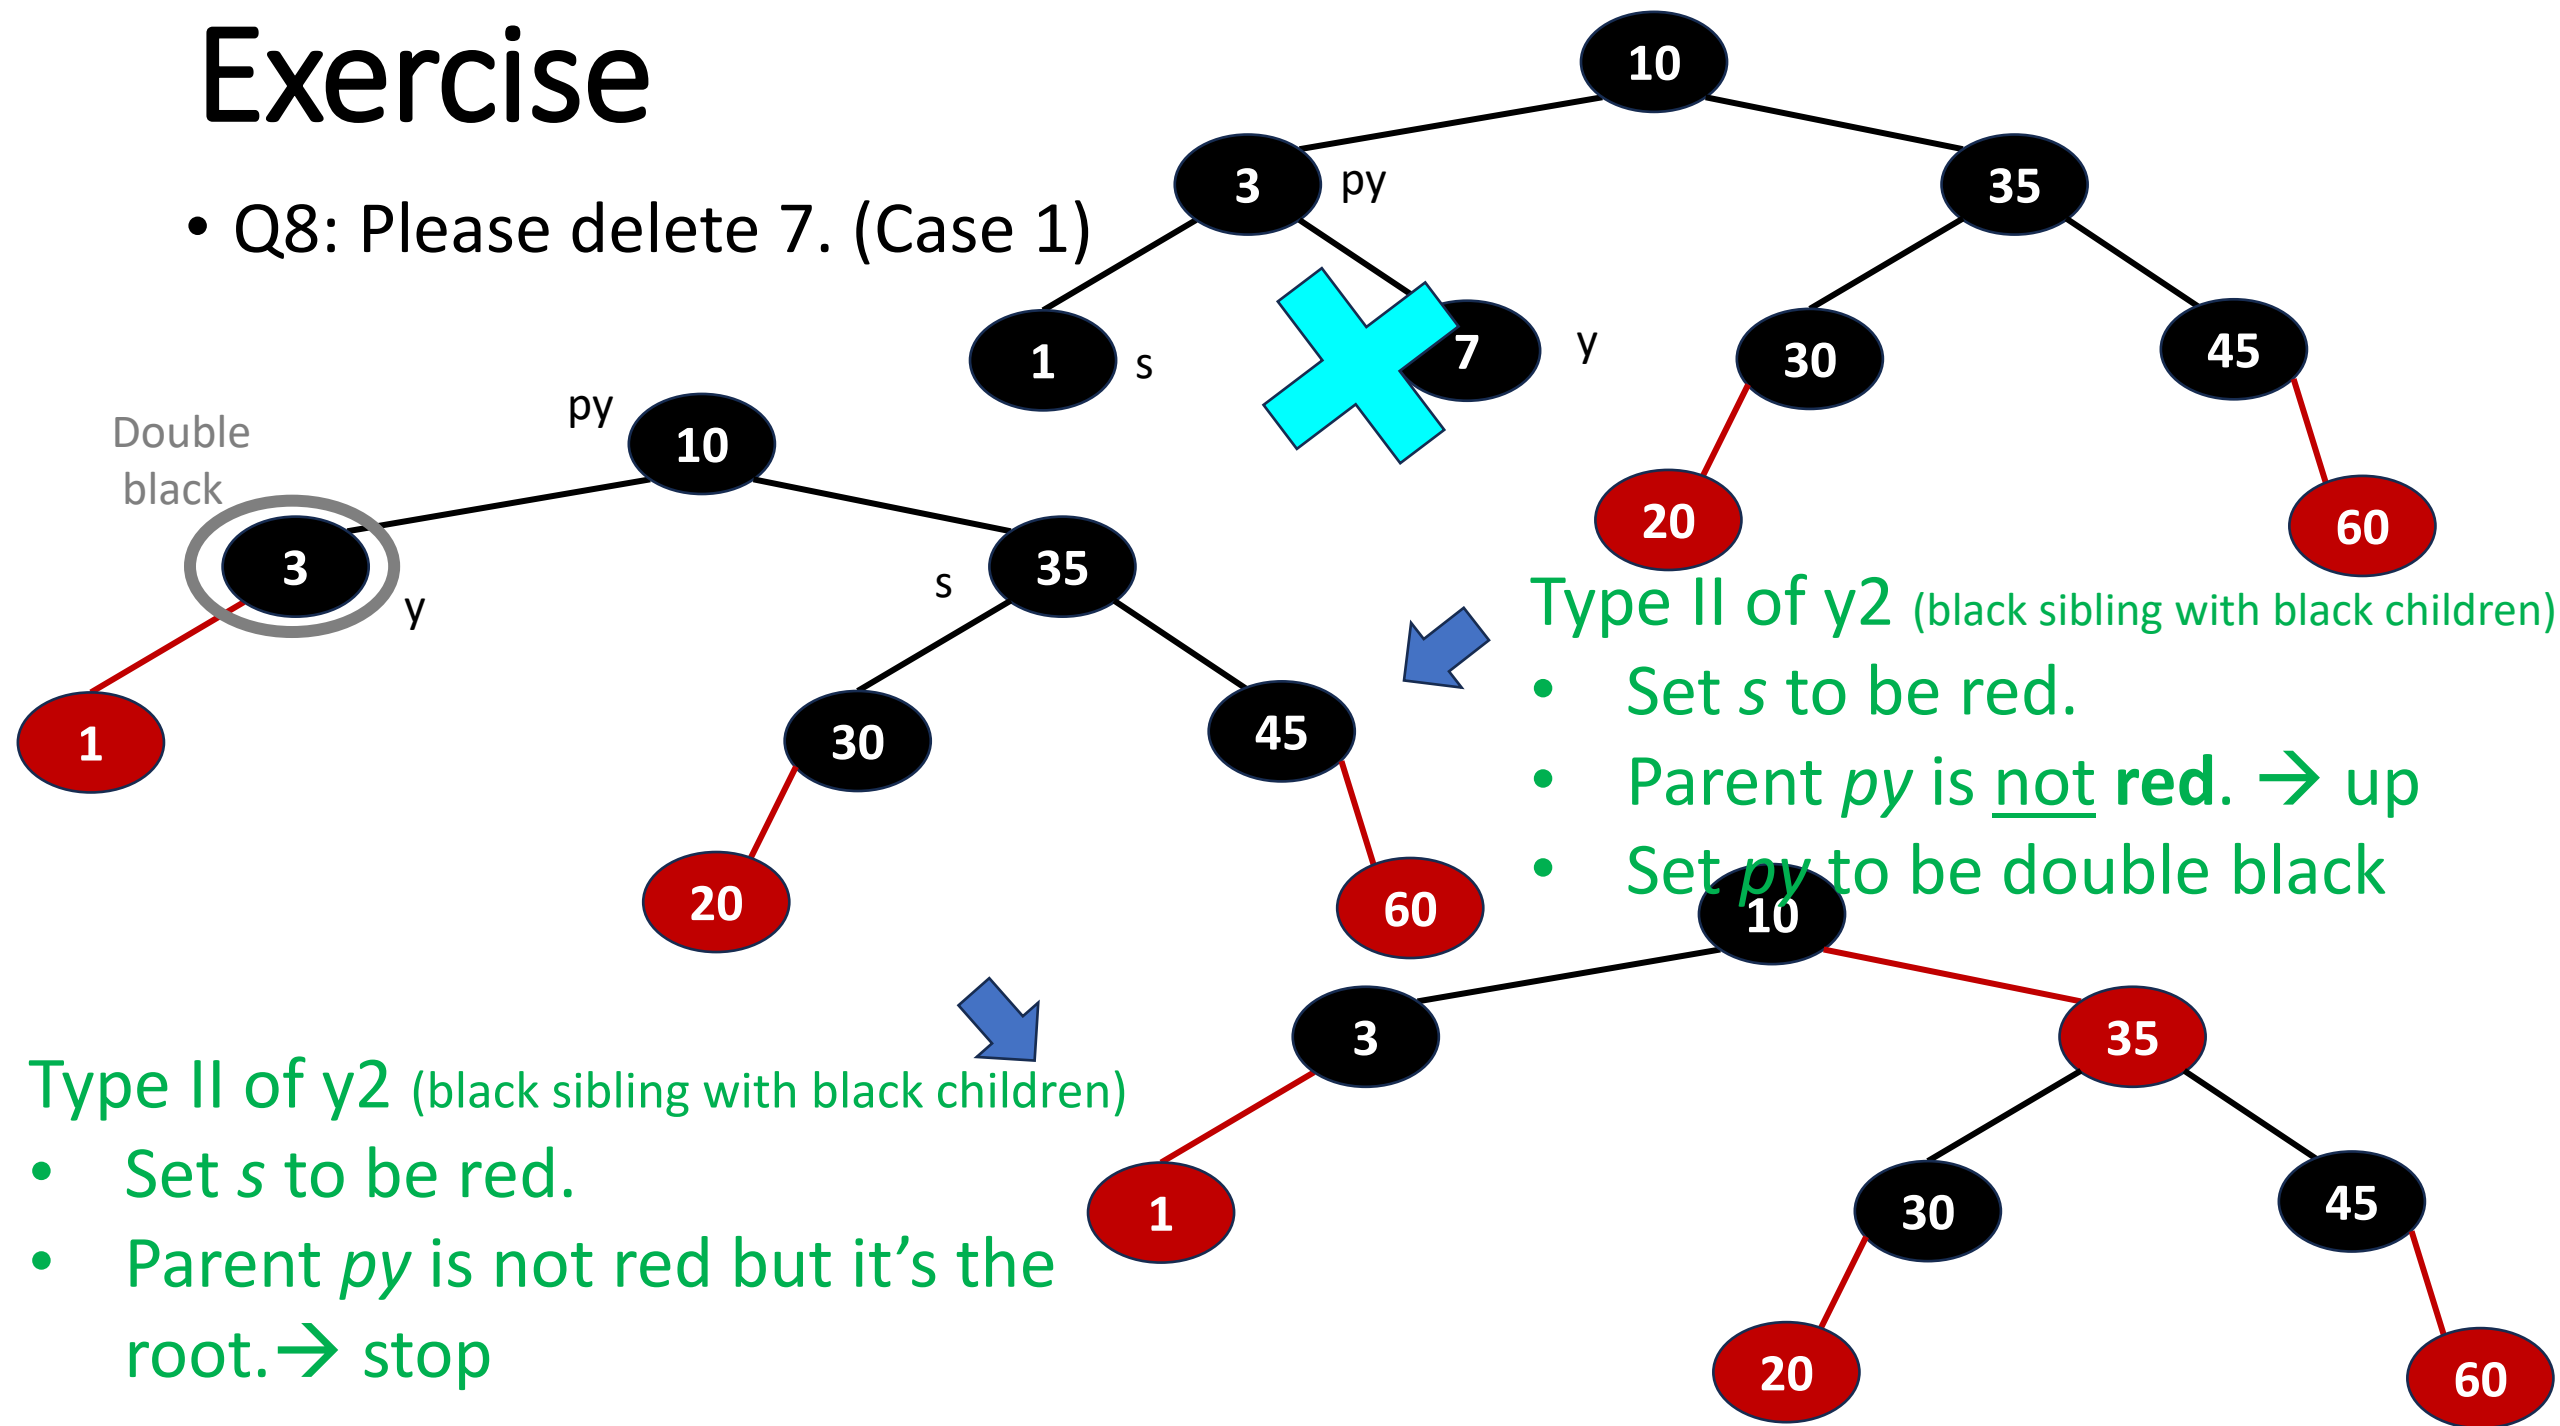
Exercise

- Q7: Please delete 40. (Case 1)



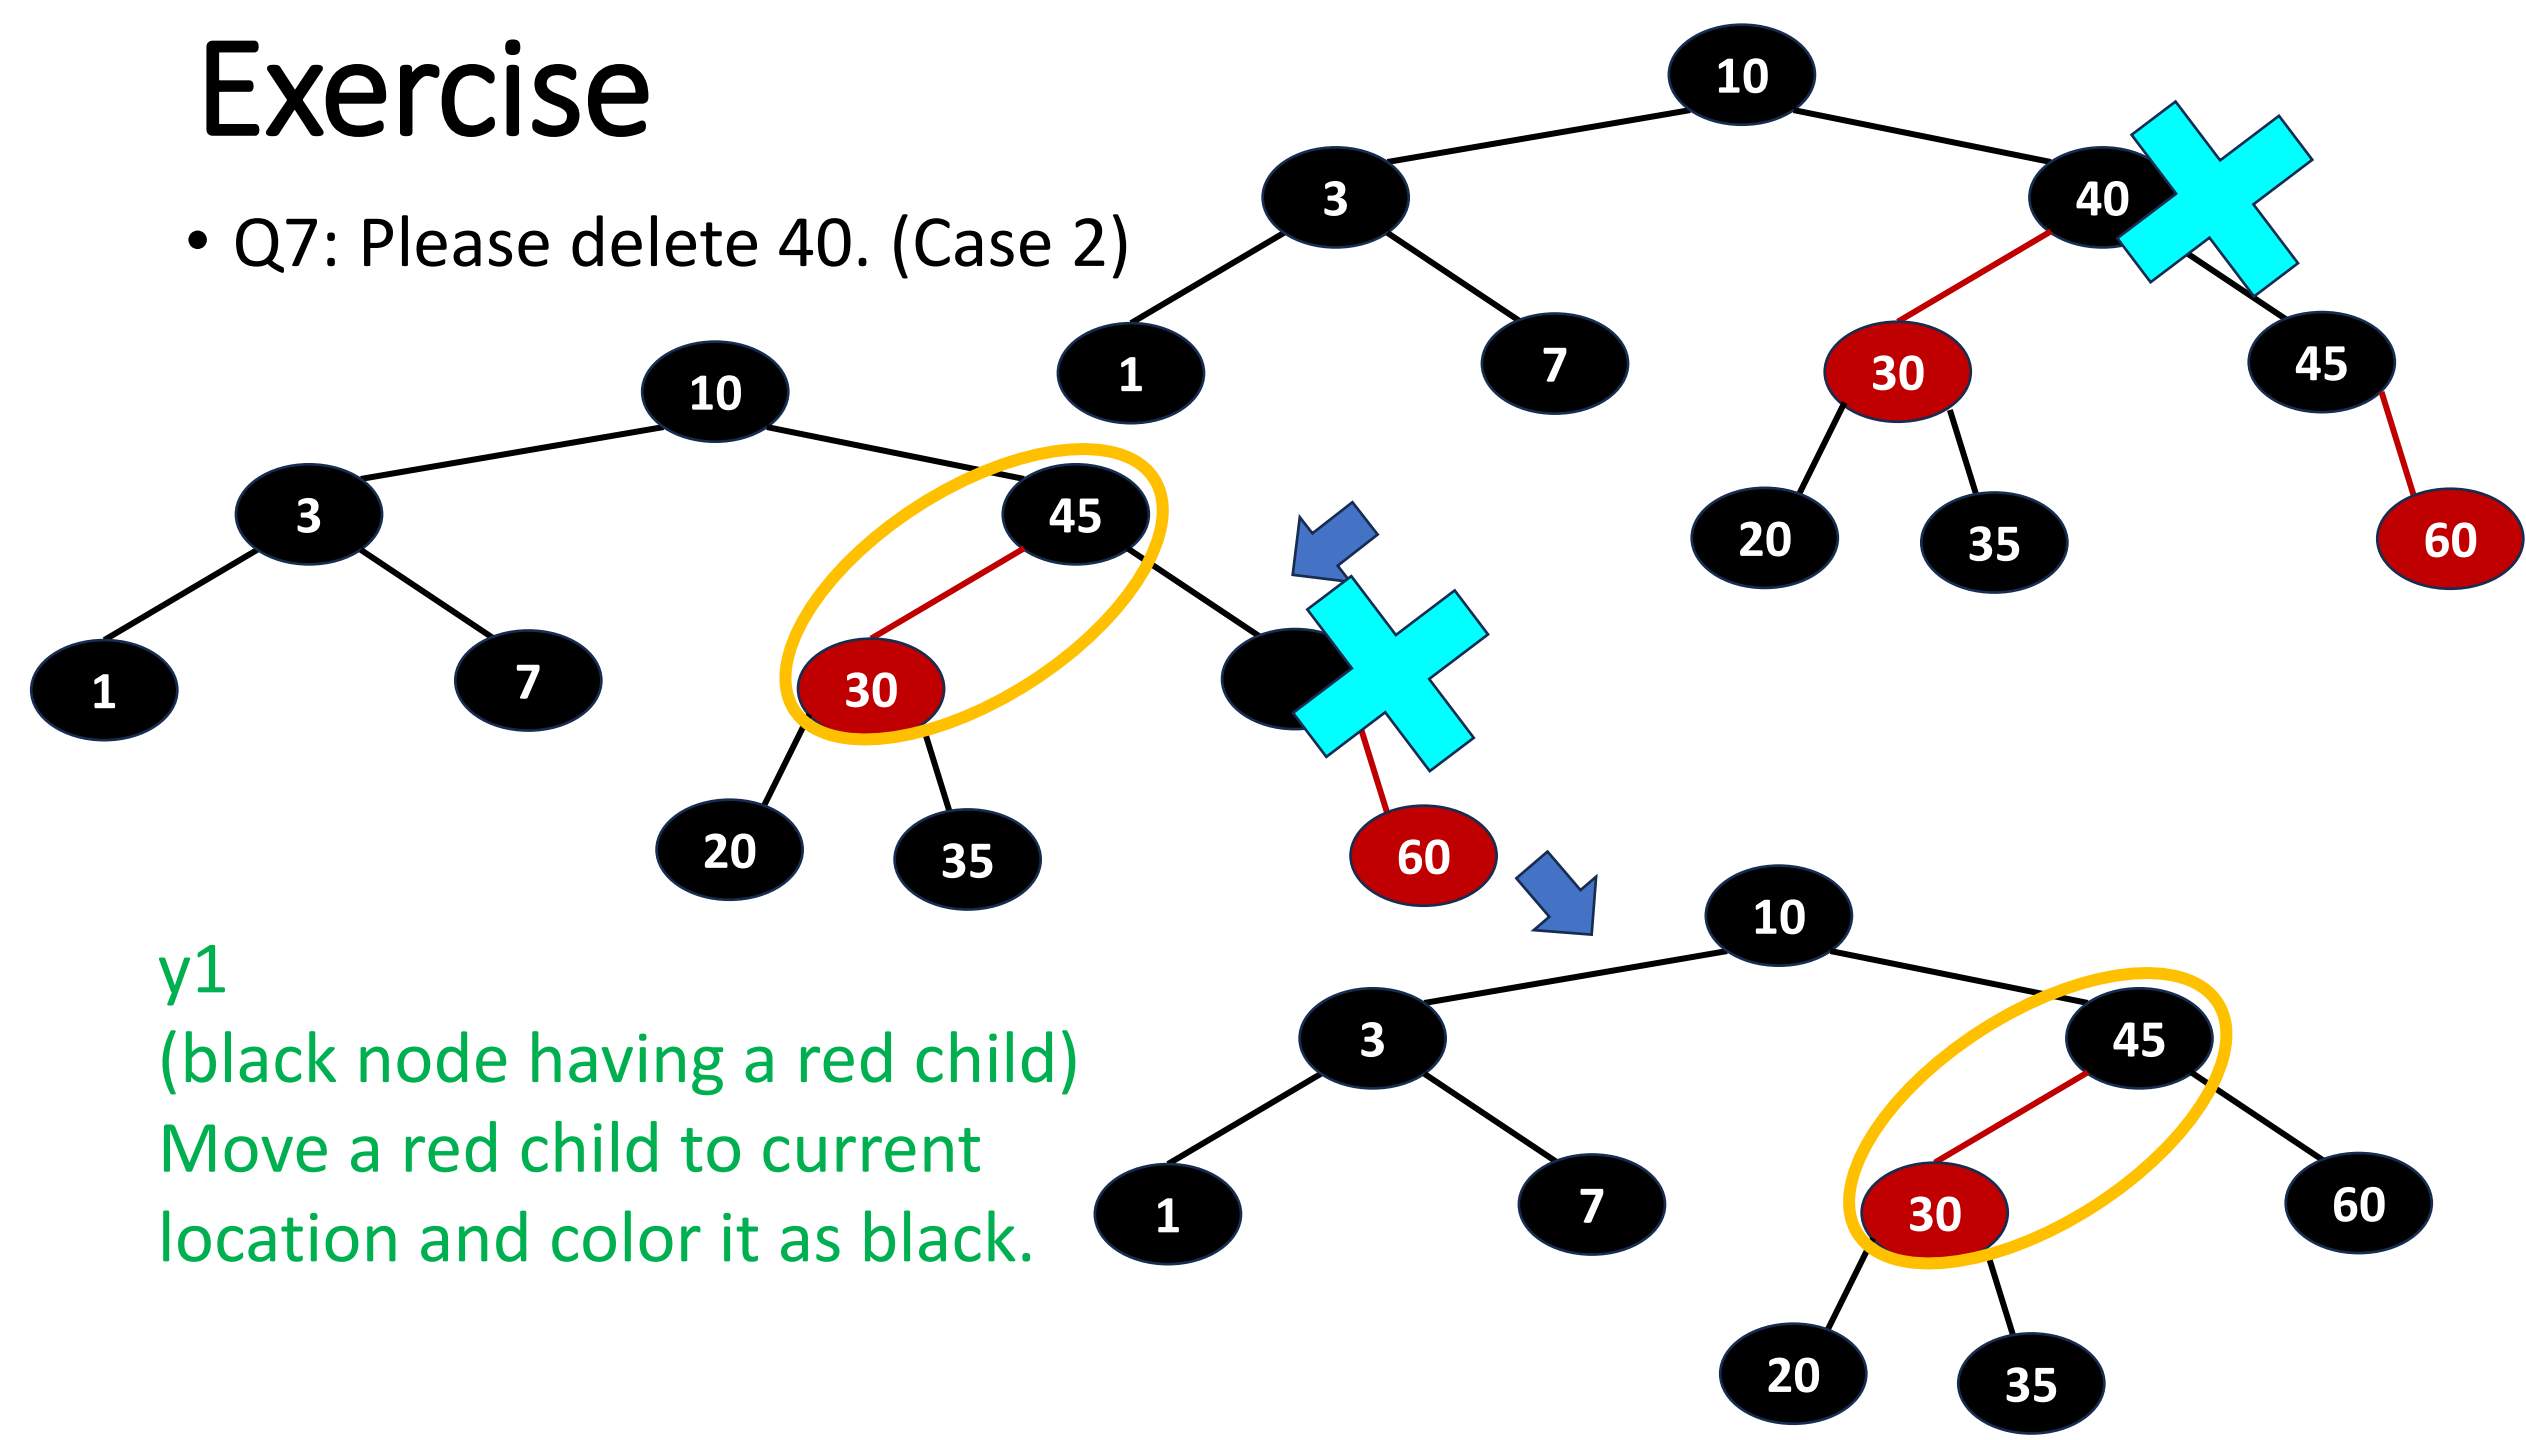
Exercise

- Q8: Please delete 7. (Case 1)



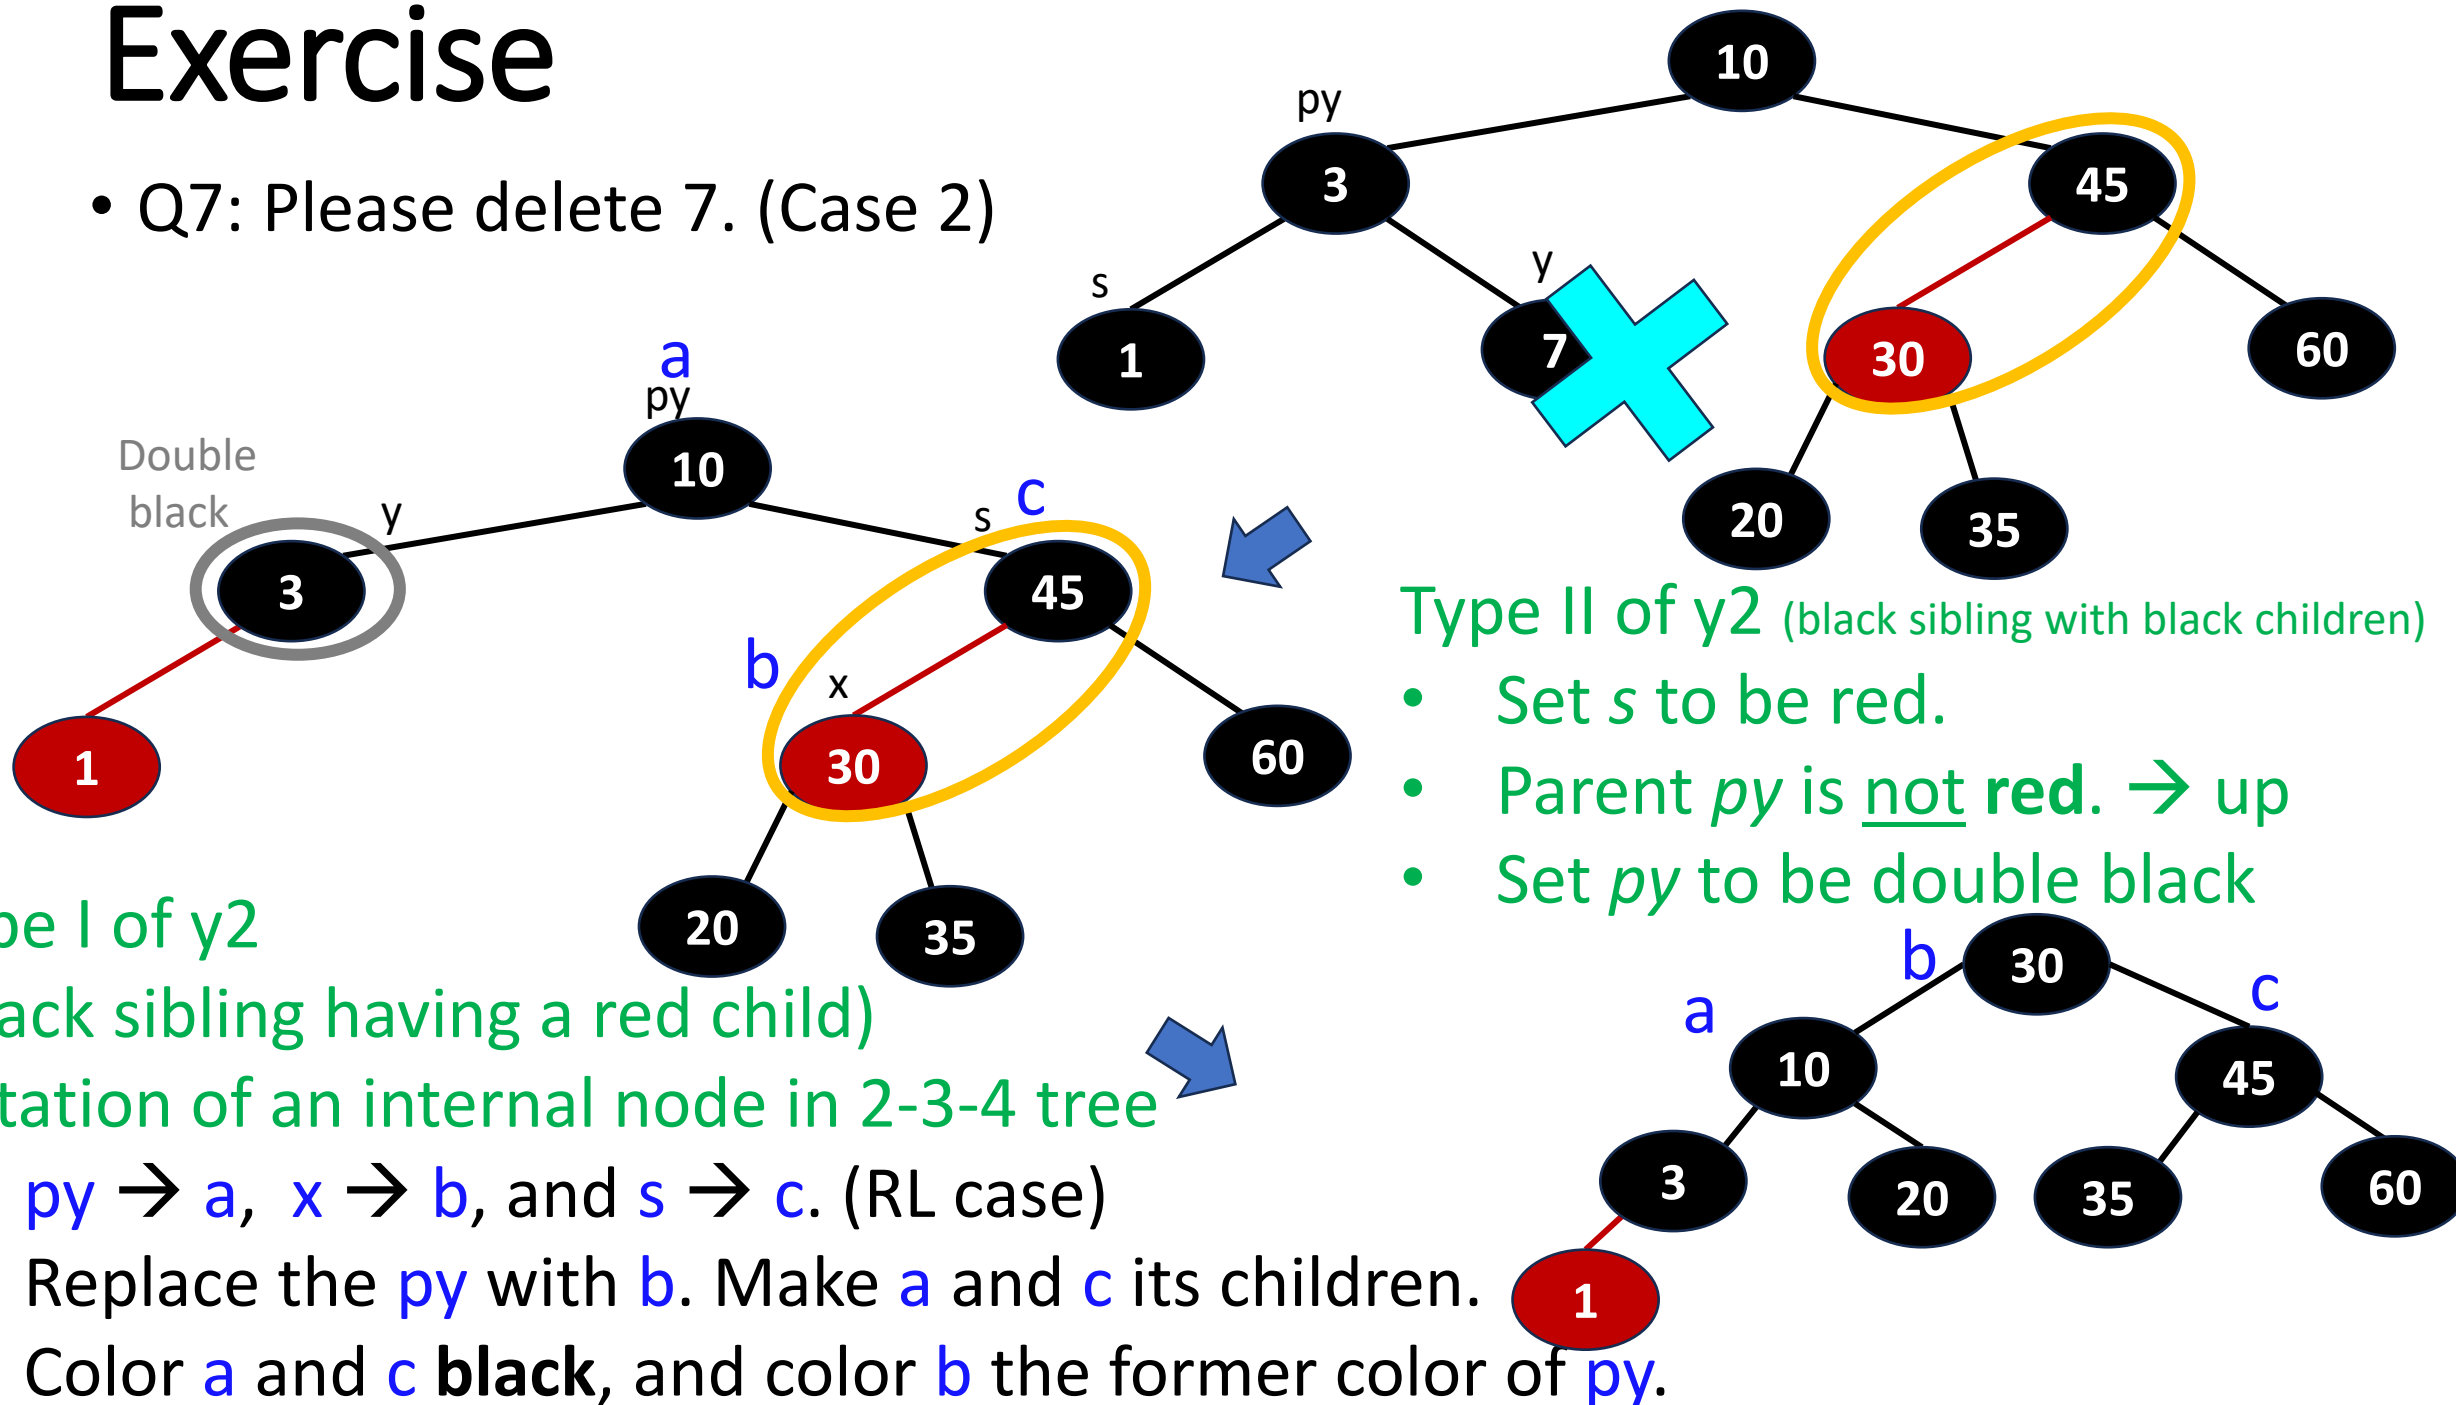
Exercise

- Q7: Please delete 40. (Case 2)



Exercise


- Q7: Please delete 7. (Case 2)



Insertion and deletion algorithms for Red-Black Trees. Are they unique?

- No!!!
- Given an 2-3-4 tree, several variant red-black trees exist

3-node



k -way search tree ($k > 4$) \rightarrow red-black tree

Main Concepts:

1. Transform a k -way search tree to 2-3-4 tree
2. Transform the 2-3-4 tree to red-black tree

Details:


1. Given a node p containing # of data pairs > 4
2. Split the keys in p into 3 groups, denoted by X , Y , and Z . Note that the # of data pairs in X , Y , or Z may be still larger than 4.
3. Consider the node p as a 4-node in 2-3-4 tree and convert X , Y , Z to three nodes in red-black tree.
4. Repeat Step 1 to 3 to build a level- i red-black tree (initially, $i=1$)
5. $i++$
6. Process the nodes containing # of data pairs larger than 4 using steps 1 to 4, and generate level- i red-black tree. ($i=2$)
7. Repeat steps 1 to 6 until the tree has only 2-, 3-, and 4-nodes. Then we convert the 2-3-4 tree to binary tree.

Note: The red-black tree will have more than two types of colors.


k -way search tree ($k > 4$) \rightarrow red-black tree

Example: $k=11$

Convert to
a 4-node




1-st level red black tree



k -way search tree ($k > 4$) \rightarrow red-black tree

2-nd level red black tree



Convert to red-black tree using the algorithm in textbook (2-3-4 tree to red-black tree)

Summary

- Red-black tree
- The relationship to 2-3-4 tree
- Operations:
 - Insertion
 - Deletion